

Operating Systems

[3. Process]

Chung-Wei Lin

cwlin@csie.ntu.edu.tw

CSIE Department

National Taiwan University

Objectives

- ❑ Identify the separate components of a process and illustrate how they are represented and scheduled in an operating system
- ❑ Describe how processes are created and terminated in an operating system, including developing programs using the appropriate system calls that perform these operations
- ❑ Describe and contrast interprocess communication using shared memory and message passing
- ❑ Design programs that use pipes and POSIX shared memory to perform interprocess communication
- ❑ Describe client-server communication using sockets and remote procedure calls

Outline

☐ **Process Concept**

➤ The Process, Process State, Process Control Block, Threads

☐ Process Scheduling

☐ Operations on Processes

☐ Interprocess Communication

☐ IPC in Shared-Memory Systems

☐ IPC in Message-Passing Systems

☐ Examples of IPC Systems

☐ Communication in Client-Server Systems

Process Concept

❑ Process: a program in execution

- A program is a passive entity
 - A file (executable file) containing a list of instructions stored on disk
- A process is an active entity with
 - A program counter specifying the next instruction to execute
 - A set of associated resources
- A program becomes a process when an executable file is loaded into memory
 - Double-click an icon representing the executable file
 - Enter the name of the executable file on the command line
- Although two processes may be associated with the same program, each of these is a separate process
 - It is also common to have a process that spawns many processes as it runs

Memory Layout of a Process

❑ Text section

- The executable code

❑ Data section

- Global variables

❑ Heap section

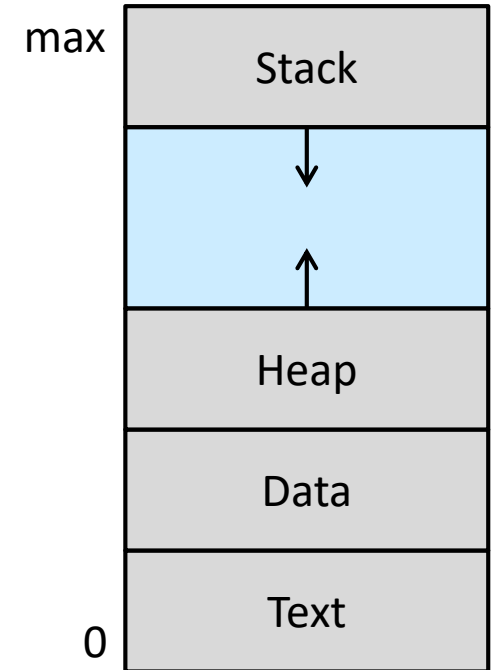
- Dynamically allocated during program run time

❑ Stack section

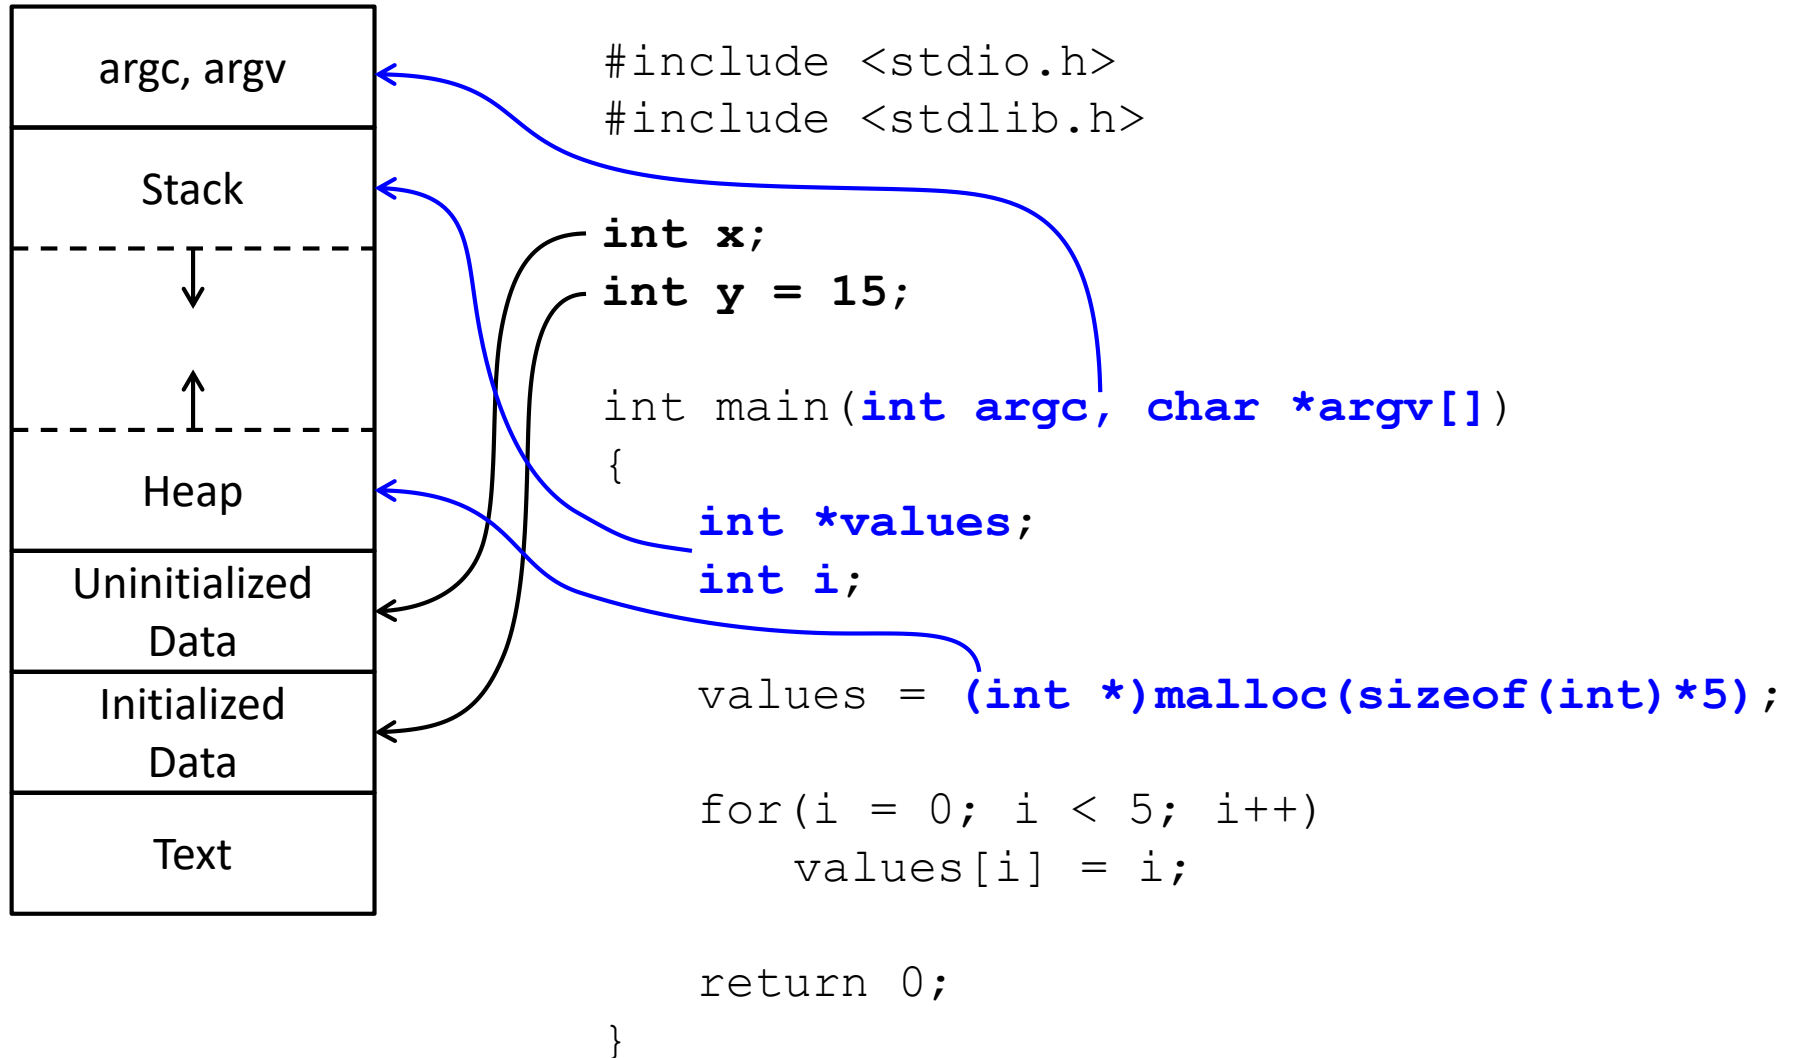
- Temporary data storage when invoking functions
- Examples: function parameters, return addresses, local variables

❑ The stack and heap sections can shrink and grow dynamically during program execution

- The operating system must ensure they do not overlap one another



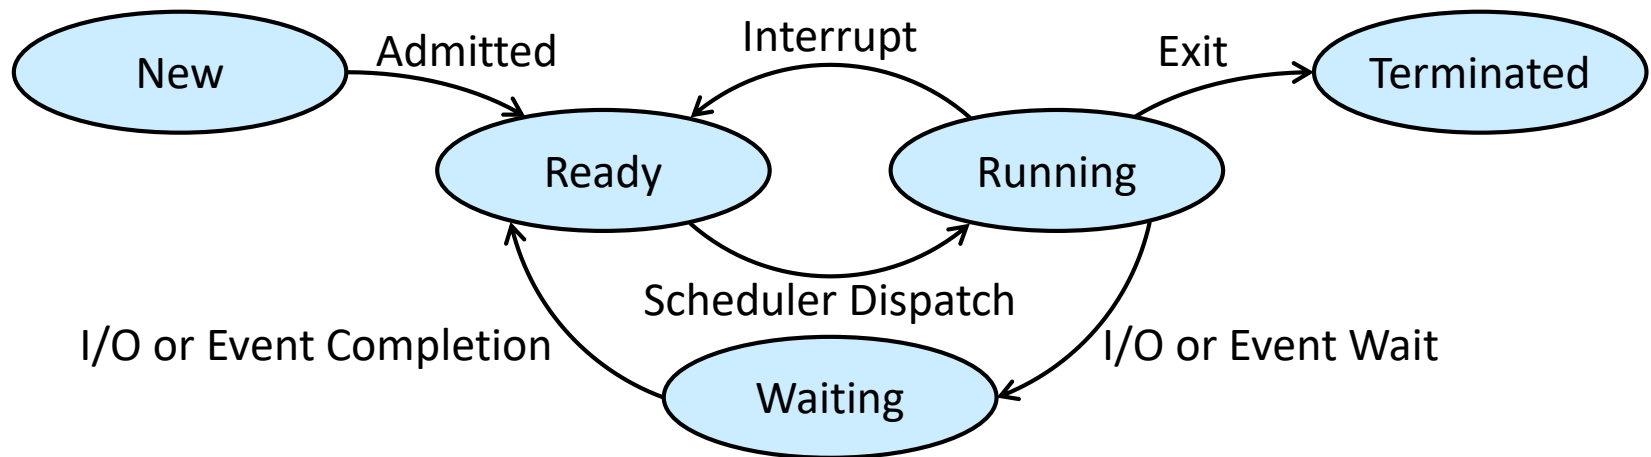
Memory Layout of a C Program



Process State

❑ As a process executes, it changes state

- **New**: the process is being created
- **Ready**: the process is waiting to be assigned to a processor
- **Running**: instructions are being executed
 - Only one process can be running on any processor core at any instant
- **Waiting**: the process is waiting for some event to occur
 - Examples: I/O completion, reception of a signal
- **Terminated**: the process has finished execution



Process Control Block

- ❑ Each process is represented in the operating system by a **process control block** (PCB), also called a **task control block**
 - Process state
 - Process number
 - Program counter
 - Address of the next instruction to be executed for this process
 - CPU registers
 - Register set where process needs to be stored for execution for running state
 - CPU-scheduling information: priority, queue pointers, etc. [Chapter 5]
 - Memory-management information: memory limits, etc. [Chapter 9]
 - Accounting information
 - Amount of CPU and real time used, time limits, account numbers, etc.
 - I/O status information
 - List of I/O devices allocated to the process, list of open files, etc.

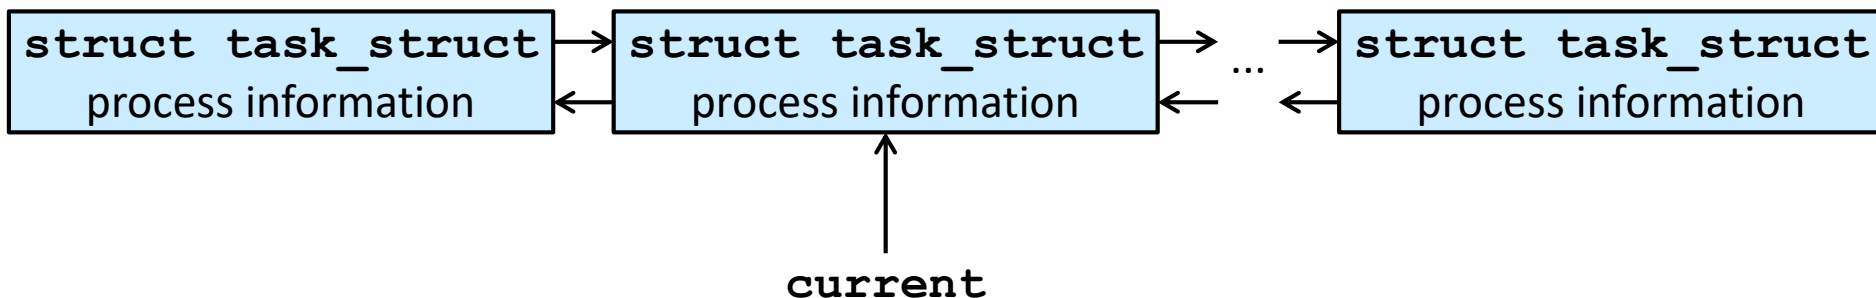
Process Representation in Linux

❑ Represented by the C structure `task_struct`

```
pid t_pid;                /* process identifier */
long state;               /* state of the process */
unsigned int time_slice   /* scheduling information */
struct task_struct *parent; /* this process's parent */
struct list_head children; /* this process's children */
struct files_struct *files; /* list of open files */
struct mm_struct *mm;      /* address space of this process */
```

❑ Within the Linux kernel, all active processes are represented using a doubly linked list of `task_struct`

- The kernel maintains a pointer `current` to the process currently executing on the system



Threads

- ❑ The process model so far has implied that a process is a program that performs a single thread of execution
 - The process performs only one task at a time
 - Example: the user cannot simultaneously type in characters and run the spell checker
- ❑ Most modern operating systems allow a process to have multiple threads of execution
 - The process performs more than one task at a time
 - Especially beneficial on multicore systems where multiple threads can run in parallel
 - The PCB is expanded to include information for each thread
 - Other changes throughout the system are also needed [Chapter 4]

Outline

- ❑ Process Concept
- ❑ **Process Scheduling**
 - Scheduling Queues, CPU Scheduling, Context Switch
- ❑ Operations on Processes
- ❑ Interprocess Communication
- ❑ IPC in Shared-Memory Systems
- ❑ IPC in Message-Passing Systems
- ❑ Examples of IPC Systems
- ❑ Communication in Client-Server Systems

Process Scheduling (1/2)

❑ Objective of multiprogramming

- To have some process running at all times and maximize CPU utilization

❑ Objective of time sharing

- To switch a CPU core among processes so frequently that users can interact with each program while it is running

Process Scheduling (2/2)

- ❑ To meet the objectives, the process scheduler selects an available process for program execution on a core
 - Each CPU core can run one process at a time
 - If there are more processes than cores, excess processes will have to wait until a core is free and can be rescheduled
 - The number of processes currently in memory is known as the degree of multiprogramming
- ❑ To meet the objectives, it also requires taking the general behavior of a process into account
 - An I/O-bound process uses more of its time doing I/O than computations
 - A CPU-bound process uses more of its time doing computations than I/O

Scheduling Queues

❑ The system includes some queues

➤ Ready queue

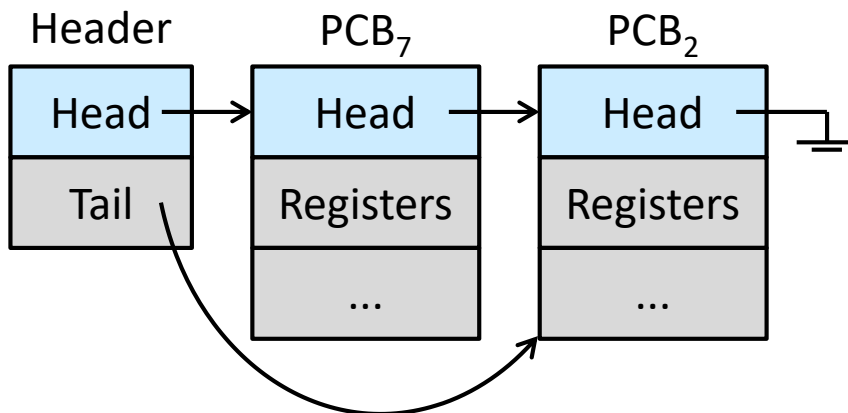
- The set of processes ready and waiting to execute on a CPU's core

➤ Wait queues

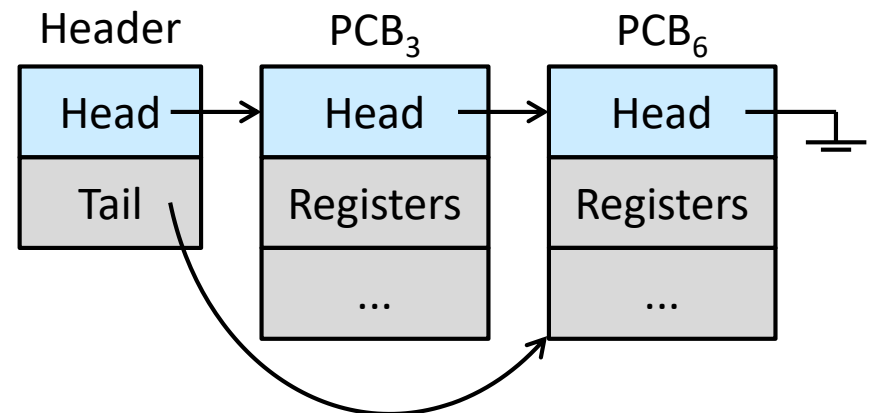
- The set of processes waiting for a certain event (e.g., completion of I/O) to occur

➤ A queue is generally stored as a linked list

Ready Queue

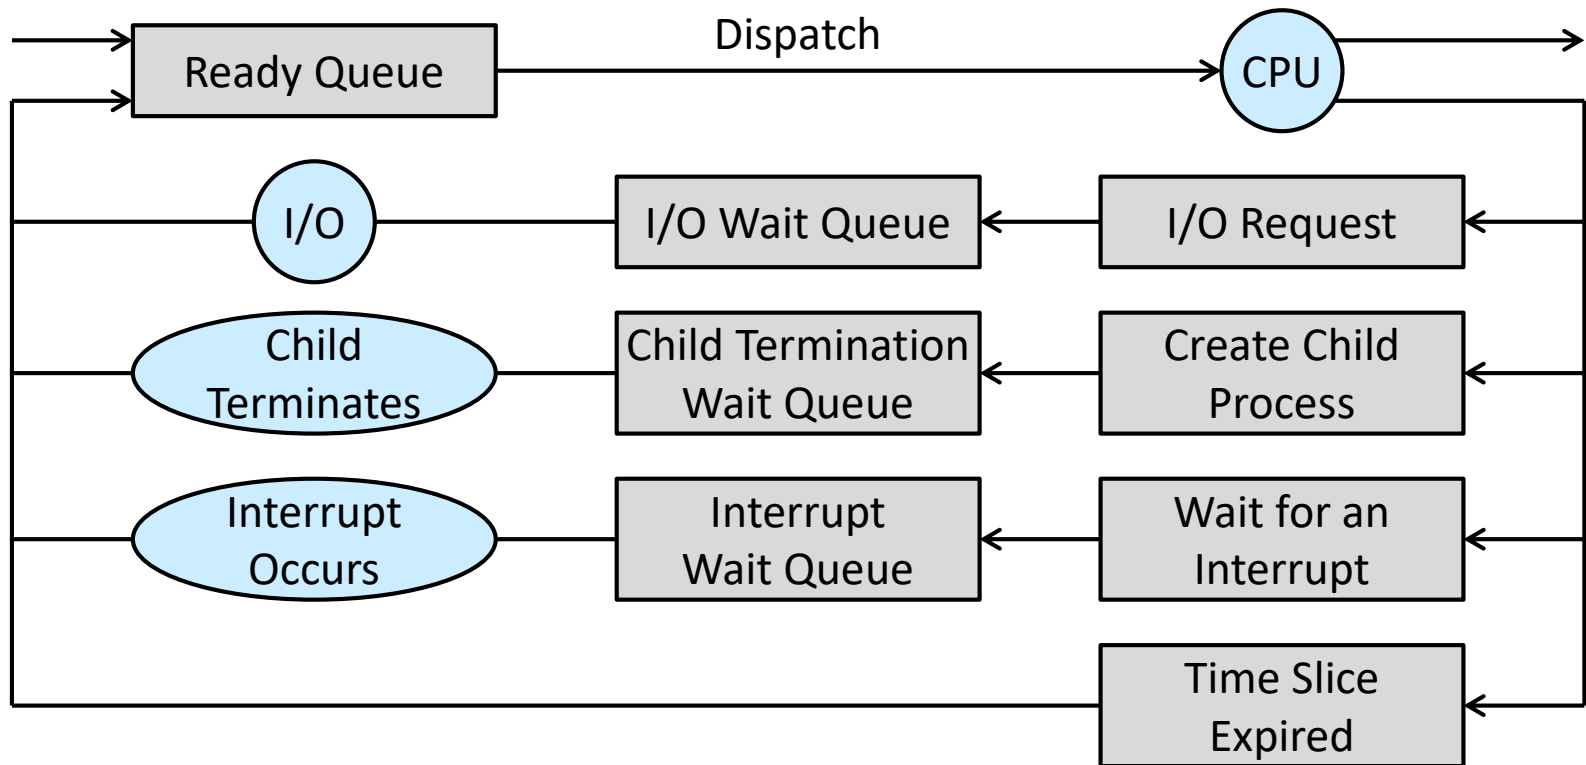


Wait Queue



Queueing Diagram

- ❑ A ready queue and a set of wait queues
- ❑ Circles represent the resources that serve the queues
- ❑ Arrows indicate the flow of processes in the system



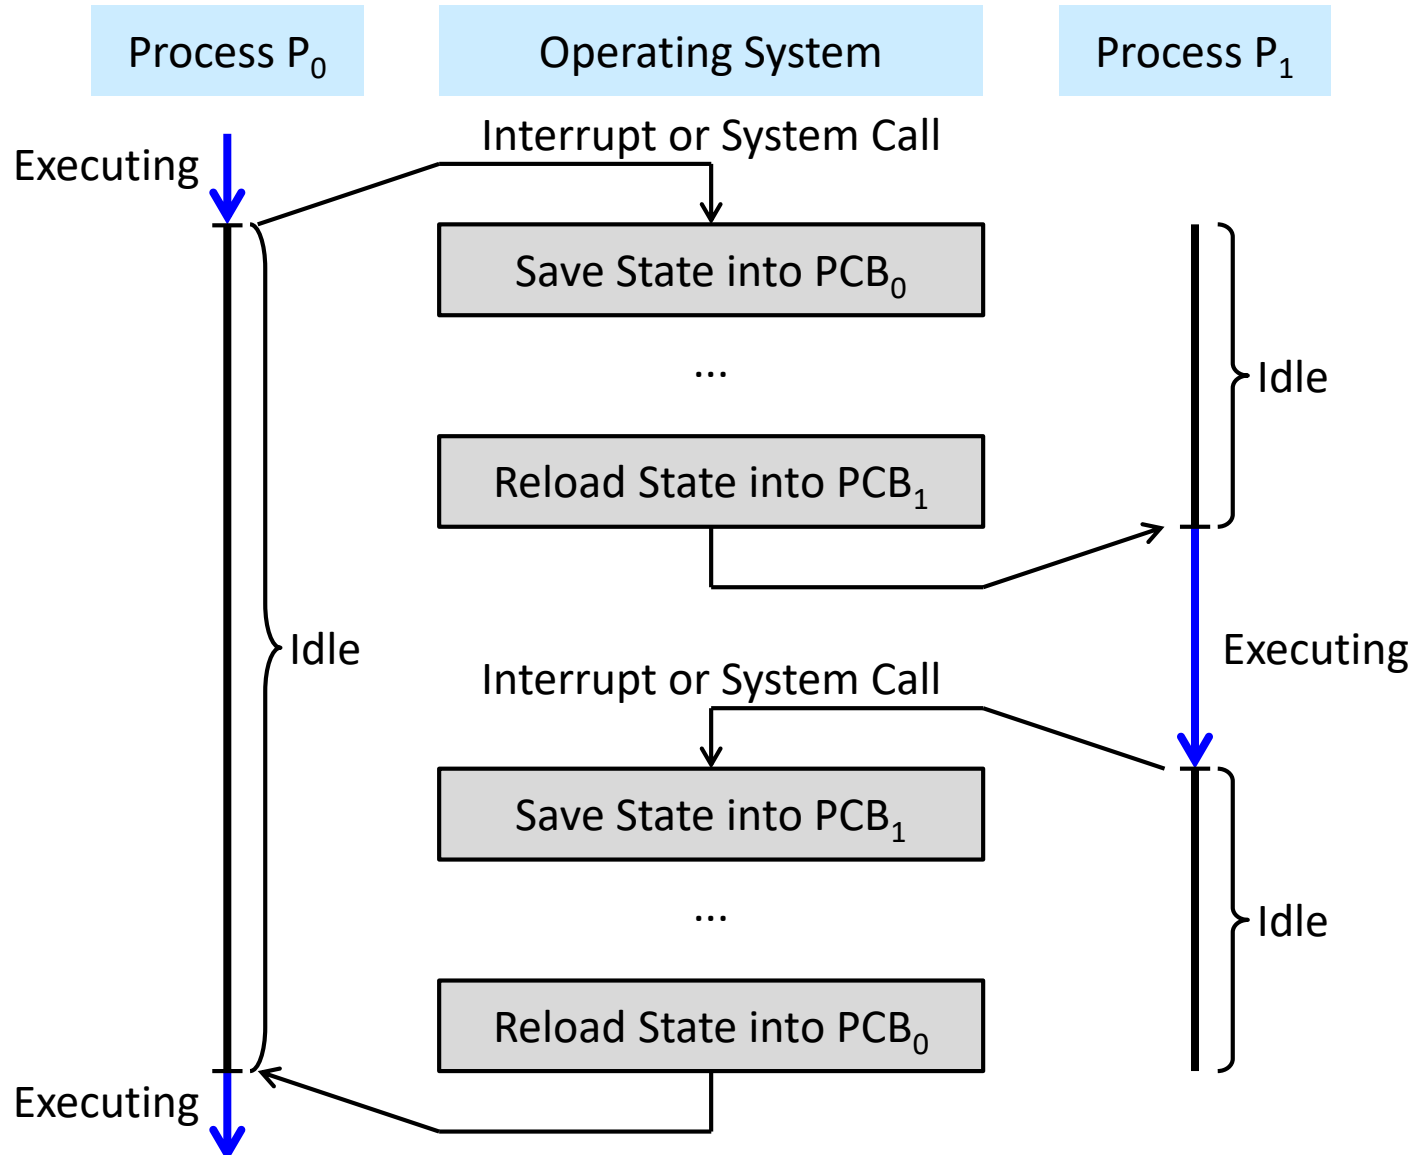
CPU Scheduling

- ❑ A process migrates among the ready queue and various wait queues throughout its lifetime
- ❑ A **CPU scheduler** executes at least once every 100 milliseconds, although typically much more frequently
 - An I/O-bound process may execute for only a few milliseconds before waiting for an I/O request
 - A CPU-bound process will require a CPU core for longer durations, but the scheduler is unlikely to grant the core to it for an extended period
- ❑ **Swapping** [Chapter 9]
 - Remove a process from memory (i.e., active contention for CPU) to disk
 - Reduce the degree of multiprogramming
 - Later, reintroduce the process into memory and continue its execution

Context Switch (1/2)

- ❑ Switching the CPU core from one process to another
- ❑ View of the current process
 - A state save of the current state of the CPU core and then a state restore to resume operations
- ❑ View of the system
 - Context switch: a state save of the current process and then a state restore of a different process
- ❑ The context of a process is represented in its PCB
 - The value of the CPU registers, the process state, and memory-management information, etc.

Diagram Showing Context Switch



Context Switch (2/2)

- ❑ Context switch time is pure overhead
 - The system does no useful work while switching
- ❑ Switching speed (usually several microseconds) depends on
 - Memory speed
 - Number of registers that must be copied
 - Existence of special instructions and hardware support
 - Example: a single instruction to load or store all registers
 - Example: multiple sets of registers, where a context switch here simply requires changing the pointer to the current register set
- ❑ The more complex the operating system, the greater the amount of work that must be done

Context Switch: Demo [Prof. Shih]

- ❑ Please check the video of the odd section
- ❑ I/O-bound process
 - Have voluntary and involuntary context switches
- ❑ CPU-bound process
 - Have involuntary context switches only or no context switch
- ❑ **`nanosleep()`** on macOS suspends the process for a very short period of time
 - Lead to involuntary context switches
- ❑ **`sleep()`** puts the process into the waiting state and the waiting queue
 - Lead to voluntary context switches

Multitasking in Mobile Systems

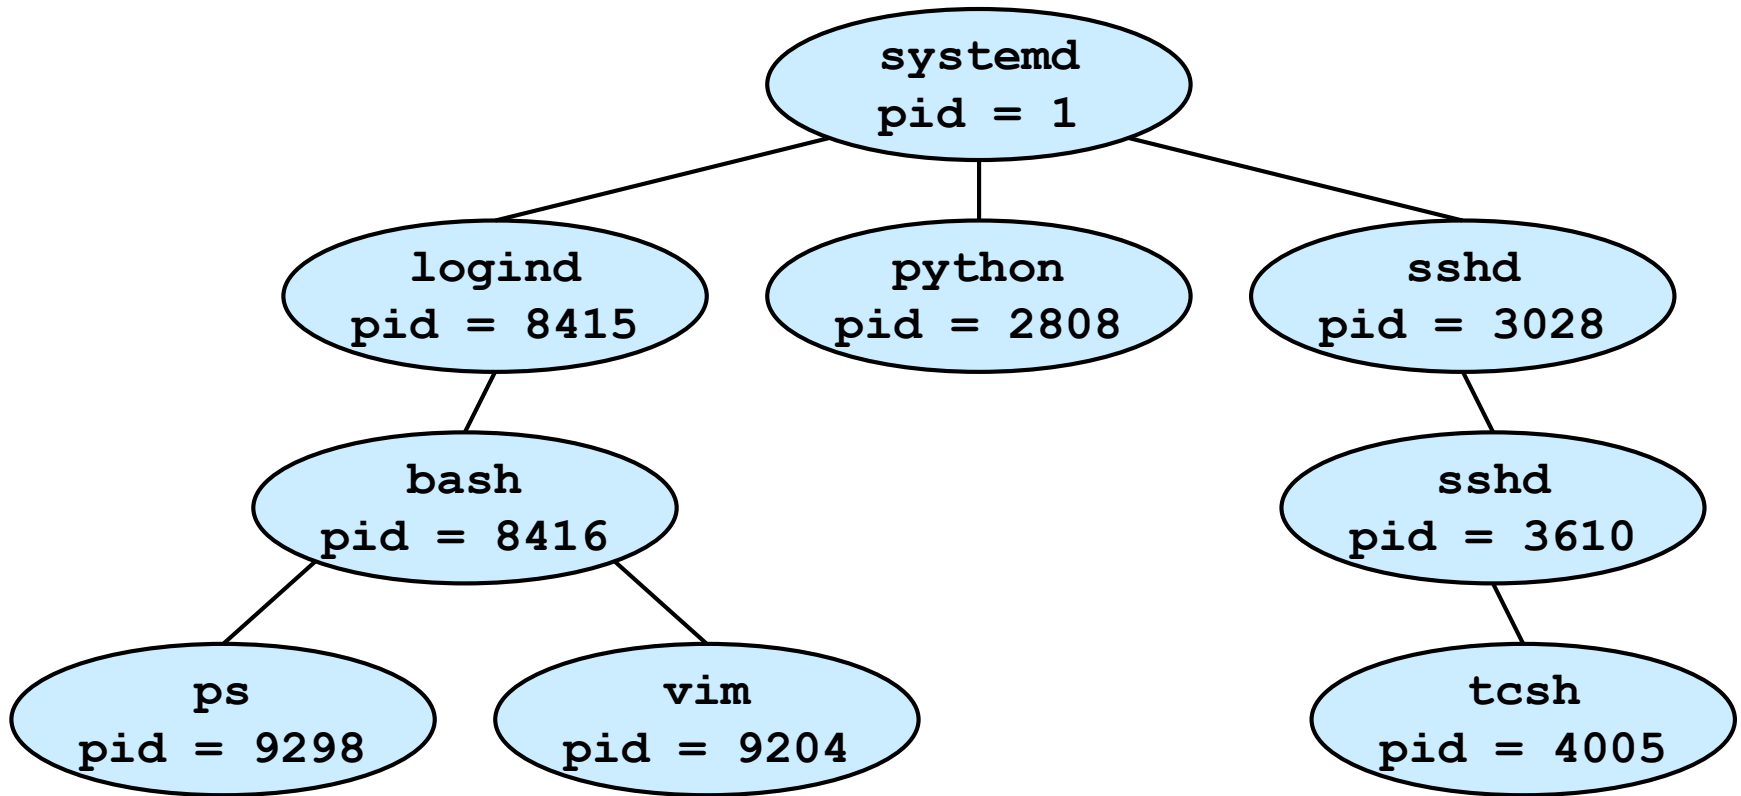
- ❑ Beginning with iOS 4, iOS allows a single **foreground** application to run with multiple **background** applications
 - Early versions of iOS does not provide user-application multitasking
 - Only one application runs in the foreground
 - All other user applications are suspended
 - **Split-screen**
 - A larger screen allows running two foreground applications at the same time
- ❑ Android runs foreground and background with fewer limits
 - A **service** runs on behalf of the background process
 - The service runs even if the background application is suspended
 - Services do not have a user interface and have a small memory footprint

Outline

- ❑ Process Concept
- ❑ Process Scheduling
- ❑ **Operations on Processes**
 - **Process Creation**, Process Termination
- ❑ Interprocess Communication
- ❑ IPC in Shared-Memory Systems
- ❑ IPC in Message-Passing Systems
- ❑ Examples of IPC Systems
- ❑ Communication in Client-Server Systems

Process Creation (1/2)

- ❑ A parent process creates children processes, which in turn create other processes, forming a tree of processes
 - Generally, processes are identified via a process identifier (pid)
- ❑ A tree of processes on a typical Linux system



Process Creation (2/2)

❑ Resource (memory, files, etc.) sharing options

- The child shares all of the parent's resources
- The child shares a subset of the parent's resources
 - Prevent any process from overloading the system by creating too many children
- The child shares none of the parent's resources

❑ Execution options

- The parent continues to execute concurrently with its children
- The parent waits until some or all of its children have terminated

❑ Address-space options

- The child is a duplicate of the parent process
 - It has the same program and data as the parent
- The child has a new program loaded into it

Process Creation Using UNIX `fork()`

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
int main() {
    pid_t pid;
    /* fork a child process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }
    return 0;
}
```

Process Creation Using UNIX `fork()`

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
int main() {
    pid_t pid;
    /* fork a child process */
    pid = fork() ;
```

❑ `fork()` creates a new process

- The child (new) process consists of a copy of the address space of the parent (original) process
- Both processes continue execution at the instruction after `fork()` with one difference
 - The return code for `fork()` is nonzero (`pid` of the child) for the parent
 - The return code for `fork()` is zero for the child

```
        printf("Child Complete");
    }
    return 0;
}
```

Process Creation Using UNIX `fork()`

❑ `exec()` loads a binary file into memory and starts execution

- `execlp()` in the example code
- It destroys the original memory image

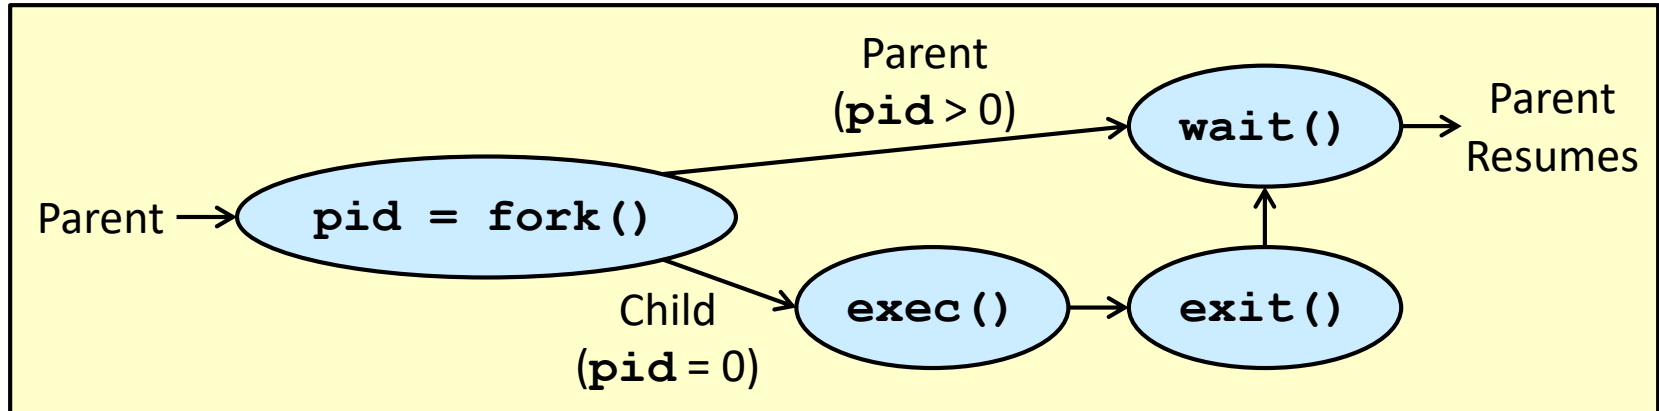
```
/* fork a child process */
pid = fork();
if (pid < 0) { /* error occurred */
    fprintf(stderr, "Fork Failed");
    return 1;
}
else if (pid == 0) { /* child process */
    execlp("/bin/ls", "ls", NULL);
}
else { /* parent process */
    /* parent will wait for the child to complete */
    wait(NULL);
    printf("Child Complete");
}
return 0;
}
```

Process Creation Using UNIX `fork()`

- ❑ The parent calls `wait()` to move itself off the ready queue until the termination of the child
 - The child terminates by either implicitly or explicitly invoking `exit()`
 - Implicitly: the C run-time library (added to UNIX executable files) includes `exit()` by default
 - Explicitly: directly use `exit()`

```
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }
    return 0;
}
```

Process Creation Using UNIX `fork()`



```
        fprintf(stderr, "Fork Failed");  
        return 1;  
    }  
    else if (pid == 0) { /* child process */  
        execlp("/bin/ls", "ls", NULL);  
    }  
    else { /* parent process */  
        /* parent will wait for the child to complete */  
        wait(NULL);  
        printf("Child Complete");  
    }  
    return 0;  
}
```

Process Creation Using Windows API

- ❑ **CreateProcess ()** requires loading a specified program into the address space of the child process at process creation
 - UNIX example: **fork ()** has the child process inheriting the address space of its parent
- ❑ **CreateProcess ()** expects no fewer than ten parameters
 - UNIX example: **fork ()** is passed no parameter
- ❑ **WaitForSingleObject ()** waits for the child process to complete
 - UNIX example: **wait ()** is equivalent
- ❑ **Note: address-space options**
 - The child is a duplicate of the parent process
 - It has the same program and data as the parent
 - The child has a new program loaded into it

Process Creation Using Windows API

```
int main(VOID){
    STARTUPINFO si;
    PROCESS_INFORMATION pi;
    /* allocate memory */
    ZeroMemory(&si, sizeof(si));
    si.cb = sizeof(si);
    ZeroMemory(&pi, sizeof(pi));
    /* create child process */
    if (!CreateProcess(NULL, /* use command line */
        "C:\\WINDOWS\\system32\\mspaint.exe", /* command */
        NULL, /* don't inherit process handle */
        NULL, /* don't inherit thread handle */
        FALSE, /* disable handle inheritance */
        0, /* no creation flags */
        NULL, /* use parent's environment block */
        NULL, /* use parent's existing directory */
        &si, &pi)){
        fprintf(stderr, "Create Process Failed");
        return -1;
    }
    /* parent will wait for the child to complete */
    WaitForSingleObject(pi.hProcess, INFINITE);
    printf("Child Complete");
    /* close handles */
    CloseHandle(pi.hProcess);
    CloseHandle(pi.hThread);
}
```

Outline

- ❑ Process Concept
- ❑ Process Scheduling
- ❑ **Operations on Processes**
 - Process Creation, **Process Termination**
- ❑ Interprocess Communication
- ❑ IPC in Shared-Memory Systems
- ❑ IPC in Message-Passing Systems
- ❑ Examples of IPC Systems
- ❑ Communication in Client-Server Systems

Process Termination (1/3)

- ❑ A process finishes executing its final statement and asks the operating system to delete it by using the **exit()** system call
 - A status value (typically an integer) is returned to its waiting parent
 - All the resources are deallocated and reclaimed by the operating system

- ❑ **exit()** provides an exit status

```
/* exit with status 1 */  
exit(1);
```

- ❑ **wait()** returns the process identifier of the terminated child

```
pid_t pid;  
int status;  
pid = wait(&status);
```

- What if there are multiple children or there is no child?
 - <http://www.csl.mtu.edu/cs4411.ck/www/NOTES/process/fork/wait.html>

Process Termination (2/3)

- ❑ A parent may terminate the execution of one of its children
 - The child has exceeded its usage of some of the resources that it has been allocated
 - The parent must have a mechanism to inspect the state of its children
 - The task assigned to the child is no longer required
 - The parent is exiting, and the operating system does not allow a child to continue if its parent terminates
 - Cascading termination: if a process terminates (either normally or abnormally), then all its children must also be terminated
 - It is normally initiated by the operating system

Process Termination (3/3)

- ❑ A **zombie** is a process that has terminated, but whose parent has not yet called **wait()**
 - When a process terminates, its resources are deallocated by the OS
 - However, its entry in the process table (containing the exit status) must remain until the parent calls **wait()**
 - Once the parent calls **wait()**, the zombie's process identifier and process-table entry are released
- ❑ An **orphan** is a process that its parent did not invoke **wait()** and terminated
 - UNIX assigns the **systemd** (or **init**) process as the new parent
 - The **systemd** (or **init**) process periodically invokes **wait()**
 - Allow the exit status of any orphan to be collected
 - Release the orphan's process identifier and process-table entry

Android Process Hierarchy

- ❑ Importance hierarchy from most to least important
 - Foreground process
 - Visible process
 - Service process
 - Background process
 - Empty process
- ❑ If system resources must be reclaimed, Android will first terminate a least-important process

Chrome Browser

❑ Multiprocess architecture: three different types of processes

- The **browser** process manages the user interface, disk and network I/O
- **Renderer** processes handles HTML, Javascript, images, and so forth
 - A new renderer process is created for each website opened in a new tab
- A **plug-in** process is created for each type of plug-in in use
 - Examples: QuickTime or Flash

❑ Each tab represents a separate process

- Better isolation
- Better security
 - Renderer processes run in a **sandbox** which means that access to disk and network I/O is restricted

Outline

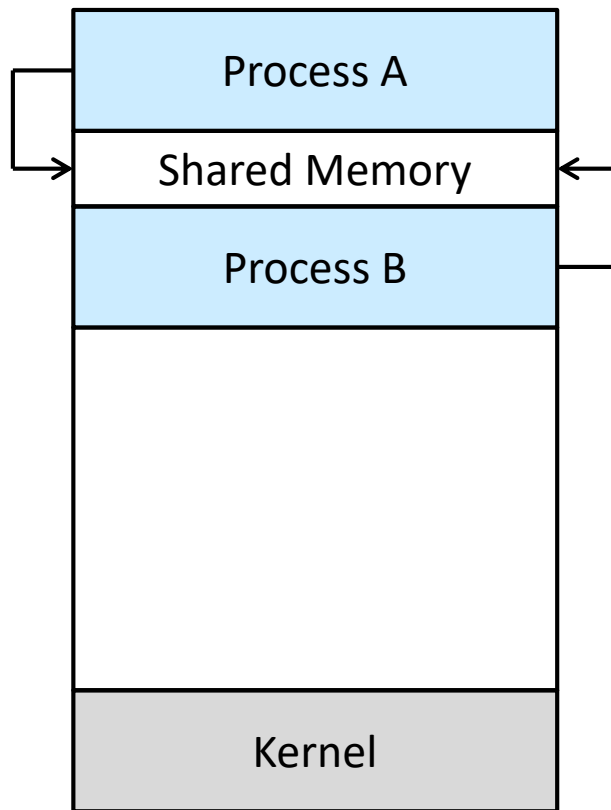
- ❑ Process Concept
- ❑ Process Scheduling
- ❑ Operations on Processes
- ❑ **Interprocess Communication**
- ❑ IPC in Shared-Memory Systems
- ❑ IPC in Message-Passing Systems
- ❑ Examples of IPC Systems
- ❑ Communication in Client-Server Systems

Interprocess Communication (1/2)

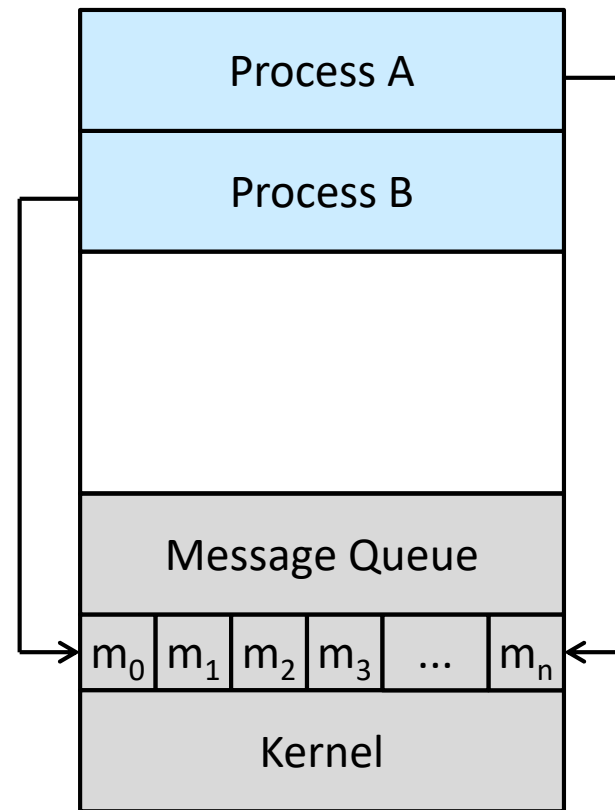
- ❑ A process may be either independent or cooperating
 - A process is **independent** if it does not share data with any other processes executing in the system
 - A process is **cooperating** if it can affect or be affected by the other processes executing in the system
 - Any process that shares data with other processes is a cooperating process
- ❑ Reasons for process cooperation
 - Information sharing, computation speedup, modularity
- ❑ Cooperating processes require an **interprocess communication** (IPC) mechanism for data exchange
 - Model 1: **shared memory**
 - Generally faster (routine memory access once shared memory is established)
 - Model 2: **message passing**
 - Easier to implement in a distributed system

Interprocess Communication (2/2)

- ❑ Cooperating processes require an interprocess communication (IPC) mechanism for data exchange



Shared Memory



Message Passing

Outline

- ❑ Process Concept
- ❑ Process Scheduling
- ❑ Operations on Processes
- ❑ Interprocess Communication
- ❑ **IPC in Shared-Memory Systems**
- ❑ IPC in Message-Passing Systems
- ❑ Examples of IPC Systems
- ❑ Communication in Client-Server Systems

IPC in Shared-Memory Systems

- ❑ Require communicating processes to establish a region of shared memory
 - Typically, the region resides in the address space of the process creating the shared-memory segment
 - Other processes attach the segment to their address space
- ❑ Normally, the operating system tries to prevent one process from accessing another process's memory
 - Shared memory requires that two or more processes agree to remove this restriction
 - The processes are also responsible for
 - The form of the data
 - The location of the data
 - Not writing to the same location simultaneously

Producer-Consumer Problem

❑ A producer process produces information that is consumed by a consumer process

➤ Examples

- A compiler may produce assembly code that is consumed by an assembler
- The assembler may produce object modules that are consumed by the loader
- A web server produces web content such as HTML files and images, which are consumed by the client web browser requesting the resource

❑ Shared memory solution

- Have a buffer of items that can be filled (produced) by the producer and emptied (consumed) by the consumer
 - This buffer resides in the shared memory
- Synchronize the producer and the consumer [Chapters 6 and 7]
 - The consumer does not try to consume an item that has not yet been produced

Types of Buffers

- ❑ The **unbounded buffer** places no practical limit on the size of the buffer
 - The producer can always produce new items
 - The consumer may have to wait for new items
- ❑ The **bounded buffer** assumes a fixed buffer size
 - The producer must wait if the buffer is full
 - The consumer must wait if the buffer is empty

Bounded Buffer

❑ The shared **buffer**

```
#define BUFFER_SIZE 10
typedef struct {
    ...
} item;
item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```

- The variable **in** points to the next free position in the buffer
- The variable **out** points to the first full position in the buffer
- The buffer is empty when **in == out**
- The buffer is full when **((in + 1) % BUFFER_SIZE) == out**
- This scheme allows at most **(BUFFER_SIZE - 1)** items in the buffer

Producer Process Using Shared Memory

❑ Producer process using shared memory

```
item next_produced;
while (true) {
    /* produce an item in next produced */
    while (((in + 1) % BUFFER_SIZE) == out)
        ; /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```

Consumer Process Using Shared Memory

❑ Consumer process using shared memory

```
item next_consumed;  
while (true) {  
    while (in == out)  
        ; /* do nothing */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    /* consume the item in next consumed */  
}
```

❑ The issue of synchronization is not addressed here

- Both of the producer process and the consumer process attempt to access the shared buffer concurrently [Chapters 6 and 7]

Outline

- ❑ Process Concept
- ❑ Process Scheduling
- ❑ Operations on Processes
- ❑ Interprocess Communication
- ❑ IPC in Shared-Memory Systems
- ❑ **IPC in Message-Passing Systems**
 - Naming, Synchronization, Buffering
- ❑ Examples of IPC Systems
- ❑ Communication in Client-Server Systems

IPC in Message-Passing Systems (1/2)

- ❑ A message-passing facility provides at least two operations
 - `send(message)`
 - `receive(message)`
- ❑ Messages can be either fixed or variable in size
 - The system-level implementation is more straightforward with fixed-sized messages
 - The programming tasks are simpler with variable-sized messages

IPC in Message-Passing Systems (2/2)

- ❑ A communication link must exist between processes sending and receiving messages
- ❑ This link can be implemented in a variety of ways
 - We are concerned here not with the link's physical implementation
 - Examples: shared memory, hardware bus, or network [Chapter 19]
 - We are concerned with its **logical implementation**
 - Direct or indirect communication
 - Synchronous or asynchronous communication
 - Automatic or explicit buffering

Outline

- ❑ Process Concept
- ❑ Process Scheduling
- ❑ Operations on Processes
- ❑ Interprocess Communication
- ❑ IPC in Shared-Memory Systems
- ❑ **IPC in Message-Passing Systems**
 - **Naming**, Synchronization, Buffering
- ❑ Examples of IPC Systems
- ❑ Communication in Client-Server Systems

Direct Communication (1/2)

- ❑ Each process that wants to communicate must explicitly name the recipient or sender of the communication

- Symmetry in addressing

- `send(P, message)`: send a message to process `P`
- `receive(Q, message)`: receive a message from process `Q`

- ❑ Properties of communication link

- A link is established automatically between every pair of processes that want to communicate
 - The processes need to know only each other's identity to communicate
- A link is associated with exactly two processes
- Between each pair of processes, there exists exactly one link

Direct Communication (2/2)

- ❑ Each process that wants to communicate must explicitly name the recipient or sender of the communication
 - Asymmetry in addressing
 - `send(P, message)`: send a message to process `P`
 - `receive(id, message)`: receive a message from any process, where `id` is set to the name of the process with which communication has taken place
- ❑ Disadvantage (both symmetric and asymmetric): limited modularity
 - Changing the identifier of a process may necessitate examining all other process definitions
 - All references to the old identifier must be found and modified to the new identifier

Indirect Communication (1/4)

- ❑ The messages are sent to and received from mailboxes (or ports)
 - **send(A, message)** : send a message to mailbox **A**
 - **receive(A, message)** : receive a message from mailbox **A**
 - A mailbox can be viewed abstractly as an object which messages can be placed into and removed from
 - Each mailbox has a unique identification
 - Two processes can communicate only if they have a shared mailbox
 - A process can communicate with another via a number of different mailboxes

Indirect Communication (2/4)

❑ Properties of indirect communication link

- A link is established between a pair of processes only if both members of the pair have a shared mailbox
- A link may be associated with more than two processes
- Between each pair of communicating processes, a number of different links may exist, with each link corresponding to one mailbox

❑ Properties of direct communication link (for comparison)

- A link is established automatically between every pair of processes that want to communicate
- A link is associated with exactly two processes
- Between each pair of processes, there exists exactly one link

Indirect Communication (3/4)

❑ Processes P_1 , P_2 , and P_3 share mailbox **A**

- Process P_1 sends a message to **A**
- Both P_2 and P_3 execute a **receive()** from **A**
- Which process will receive the message sent by P_1 ?

❑ Some solutions

- Allow a link to be associated with two processes at most
- Allow at most one process at a time to execute a **receive()** operation
- Allow the system to select arbitrarily which process will receive the message
 - Define an algorithm for selecting which process will receive the message

Indirect Communication (4/4)

❑ A mailbox owned by a process (the mailbox is part of the address space of the process)

- When a process that owns a mailbox terminates, the mailbox disappears
- Any process that subsequently sends a message to this mailbox must be notified that the mailbox no longer exists

❑ A mailbox owned by the operating system

- The operating system must allow a process to
 - Create a new mailbox
 - Send and receive messages through the mailbox
 - Delete a mailbox
- The process that creates a new mailbox is the mailbox's owner
- The ownership and receiving privilege may be passed to other processes through appropriate system calls

Outline

- ❑ Process Concept
- ❑ Process Scheduling
- ❑ Operations on Processes
- ❑ Interprocess Communication
- ❑ IPC in Shared-Memory Systems
- ❑ **IPC in Message-Passing Systems**
 - Naming, Synchronization, Buffering
- ❑ Examples of IPC Systems
- ❑ Communication in Client-Server Systems

Synchronization

❑ Message passing may be either blocking (synchronous) or nonblocking (asynchronous)

➤ Blocking send

- The sending process is blocked until the message is received by the receiving process or by the mailbox

➤ Nonblocking send

- The sending process sends the message and resumes operation

➤ Blocking receive

- The receiver blocks until a message is available

➤ Nonblocking receive

- The receiver retrieves either a valid message or a null

❑ Different combinations of **send()** and **receive()** are possible

- When both send() and receive() are blocking, we have a rendezvous

Blocking Producer and Consumer

- ❑ The solution is trivial with blocking `send()` and `receive()`
- ❑ The producer process

```
message next_produced;  
while (true) {  
    /* produce an item in next_produced */  
    send(next_produced);  
}
```

- ❑ The consumer process

```
message next_consumed;  
while (true) {  
    receive(next_consumed);  
    /* consume the item in next_consumed */  
}
```

Outline

- ❑ Process Concept
- ❑ Process Scheduling
- ❑ Operations on Processes
- ❑ Interprocess Communication
- ❑ IPC in Shared-Memory Systems
- ❑ **IPC in Message-Passing Systems**
 - Naming, Synchronization, **Buffering**
- ❑ Examples of IPC Systems
- ❑ Communication in Client-Server Systems

Buffering

- ❑ Messages exchanged by communicating processes reside in a temporary queue
 - No matter communication is direct or indirect
- ❑ Basically, such queues can be implemented in three ways
 - Zero capacity
 - The sender must block until the recipient receives the message
 - Bounded capacity
 - If the link is full, the sender must block until space is available in the queue
 - Unbounded capacity
 - The sender never blocks

Outline

- ❑ Process Concept
- ❑ Process Scheduling
- ❑ Operations on Processes
- ❑ Interprocess Communication
- ❑ IPC in Shared-Memory Systems
- ❑ IPC in Message-Passing Systems
- ❑ **Examples of IPC Systems**
 - **POSIX Shared Memory**, Mach Message Passing, Windows, Pipes
- ❑ Communication in Client-Server Systems

POSIX Shared Memory

❑ Portable Operating System Interface (POSIX)

- A family of standards specified by the IEEE Computer Society for maintaining compatibility between operating systems [Wikipedia]
 - Variants of Unix and other operating systems

❑ Several IPC mechanisms are available for POSIX systems

❑ Here, we explore the POSIX API for shared memory

- Create a shared-memory object
 - `shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);`
- Configure the size of the object in bytes
 - `ftruncate(shm_fd, 4096);`
- Establish a memory-mapped file containing the shared-memory object and return a pointer to the memory-mapped file
 - `mmap()`

https://man7.org/linux/man-pages/man3/shm_open.3.html

<https://man7.org/linux/man-pages/man2/mmap.2.html>

Producer with POSIX Shared Memory

```
int main() {
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* strings written to shared memory */
    const char *message_0 = "Hello";
    const char *message_1 = "World!";
    /* shared memory file descriptor */
    int fd;
    /* pointer to shared memory object */
    char *ptr;
    /* create the shared memory object */
    fd = shm_open(name, O_CREAT | O_RDWR, 0666);
    /* configure the size of the shared memory object */
    ftruncate(fd, SIZE);
    /* memory map the shared memory object */
    ptr = (char *) mmap(0, SIZE, PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0);
    /* write to the shared memory object */
    sprintf(ptr, "%s", message_0);
    ptr += strlen(message_0);
    sprintf(ptr, "%s", message_1);
    ptr += strlen(message_1);
    return 0;
}
```

Consumer with POSIX Shared Memory

```
int main() {
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* shared memory file descriptor */
    int fd;
    /* pointer to shared memory object */
    char *ptr;
    /* open the shared memory object */
    fd = shm_open(name, O_RDONLY, 0666);
    /* memory map the shared memory object */
    ptr = (char *) mmap(0, SIZE, PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0);
    /* read from the shared memory object */
    printf("%s", (char *)ptr);
    /* remove the shared memory object */
    shm_unlink(name);
    return 0;
}
```

Outline

- ❑ Process Concept
- ❑ Process Scheduling
- ❑ Operations on Processes
- ❑ Interprocess Communication
- ❑ IPC in Shared-Memory Systems
- ❑ IPC in Message-Passing Systems
- ❑ **Examples of IPC Systems**
 - POSIX Shared Memory, **Mach Message Passing**, Windows, Pipes
- ❑ Communication in Client-Server Systems

Mach Message Passing (1/3)

- ❑ Mach was included in the macOS and iOS operating systems
 - Mach was initially designed for distributed systems
- ❑ Most communication in Mach, including all inter-task communication, is carried out by messages
 - Create a new **port** (mailbox) and allocate space for its message queue
 - `mach_port_allocate()`
 - Send and receive messages
 - `mach_msg()`
- ❑ Two special ports are created with a task (similar to a process)
 - The kernel has receive rights to the **Task Self** port
 - The task can send messages to the kernel
 - The task has receive rights to the **Notify** port
 - The kernel can send notification of event occurrences to the task

Mach Message Passing (2/3)

❑ Two message fields

- A fixed-size message header and a variable-sized body

```
struct message {  
    mach_msg_header_t header;  
    int data;  
};
```

❑ The send and receive operations themselves are flexible

- Example: when a message is sent to a port, if the port's queue is full, the sender has several options (specified via parameters to `mach_msg()`)
 - Wait indefinitely until there is room in the queue
 - Wait at most *n* milliseconds
 - Do not wait at all but rather return immediately
 - Temporarily cache a message

Mach Message Passing (3/3)

```
mach_port_t client;
mach_port_t server;

/* Client Code */
struct message message;
// construct the header
message.header.msgh_size = sizeof(message);
message.header.msgh_remote_port = server;
message.header.msgh_local_port = client;
// send the message
mach_msg(&message.header, // message header
        MACH_SEND_MSG, // sending a message
        sizeof(message), // size of message sent
        0, // maximum size of received message - unnecessary
        MACH_PORT_NULL, // name of receive port - unnecessary
        MACH_MSG_TIMEOUT_NONE, // no time outs
        MACH_PORT_NULL // no notify port
);

/* Server Code */
struct message message;
// receive the message
mach_msg(&message.header, // message header
        MACH_RCV_MSG, // sending a message --> receiving a message
        0, // size of message sent
        sizeof(message), // maximum size of received message
        server, // name of receive port
        MACH_MSG_TIMEOUT_NONE, // no time outs
        MACH_PORT_NULL // no notify port
);
```

Outline

- ❑ Process Concept
- ❑ Process Scheduling
- ❑ Operations on Processes
- ❑ Interprocess Communication
- ❑ IPC in Shared-Memory Systems
- ❑ IPC in Message-Passing Systems
- ❑ **Examples of IPC Systems**
 - POSIX Shared Memory, Mach Message Passing, **Windows**, Pipes
- ❑ Communication in Client-Server Systems

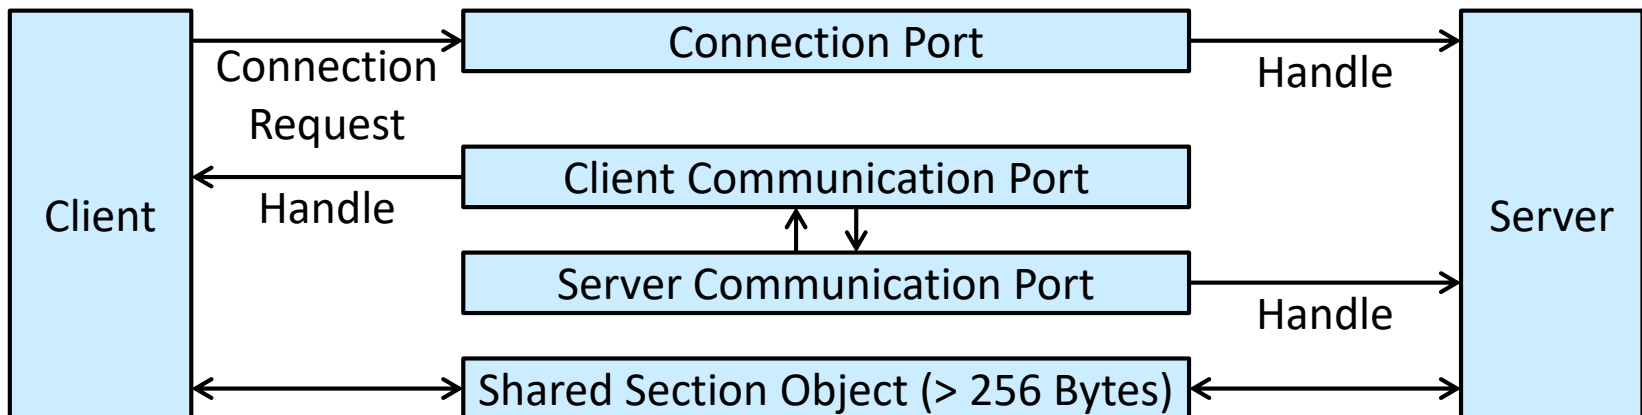
Windows

❑ Support multiple operating environments (subsystems)

- Application programs communicate with these subsystems via a message-passing mechanism
- Application programs can be considered clients of a subsystem server

❑ Advanced local procedure call (ALPC) facility

- The client opens a handle (abstract reference) to the server's **connection-port** object and sends a connection request to that port
- The server creates a channel and returns a handle to the client
 - The channel consists of two private communication ports for client-server messages and for server-client messages



Outline

- ❑ Process Concept
- ❑ Process Scheduling
- ❑ Operations on Processes
- ❑ Interprocess Communication
- ❑ IPC in Shared-Memory Systems
- ❑ IPC in Message-Passing Systems
- ❑ **Examples of IPC Systems**
 - POSIX Shared Memory, Mach Message Passing, Windows, **Pipes**
- ❑ Communication in Client-Server Systems

Pipes

❑ A conduit allowing two processes to communicate

- Pipes were one of the first IPC mechanisms in early UNIX systems

❑ Simple but with some issues

- Bidirectional or unidirectional?
- If bidirectional, half duplex (only one direction at a time) or full duplex (both directions at a time)?
- Must a relationship (such as parent-child) exist between the communicating processes?
- Can the pipes communicate over a network, or must the communicating processes reside on the same machine?

❑ Two common types used on both UNIX and Windows systems

- Ordinary pipes
- Named pipes

Ordinary Pipes

- ❑ Ordinary pipes allow two processes to communicate in standard producer-consumer fashion
 - The producer writes to one end of the pipe (the write end)
 - The consumer reads from the other end (the read end)
- ❑ An ordinary pipe is unidirectional
- ❑ An ordinary pipe cannot be accessed from outside the process that created it
- ❑ Ordinary pipes on Windows are termed anonymous pipes

Ordinary Pipes in UNIX

```
#define BUFFER_SIZE 25
#define READ_END 0
#define WRITE_END 1
int main(void) {
    char write_msg[BUFFER_SIZE] = "Greetings";
    char read_msg[BUFFER_SIZE];
    int fd[2];
    pid_t pid;
    pipe(fd); /* create the pipe */
    /* omit codes if pipe failed */
    pid = fork(); /* fork a child process */
    /* omit codes if fork failed */
    if (pid > 0) { /* parent process */
        close(fd[READ_END]); /* close the unused end of the pipe */
        write(fd[WRITE_END], write_msg, strlen(write_msg)+1); /* write to the pipe */
        close(fd[WRITE_END]); /* close the write end of the pipe */
    }
    else { /* child process */
        close(fd[WRITE_END]); /* close the unused end of the pipe */
        read(fd[READ_END], read_msg, BUFFER_SIZE); /* read from the pipe */
        printf("read %s", read_msg);
        close(fd[READ_END]); /* close the read end of the pipe */
    }
    return 0;
}
```

Named Pipes

- ❑ Named pipes provide a much more powerful communication tool
 - Communication can be bidirectional
 - No parent-child relationship is required
 - Several processes can use a named pipe for communication
 - In a typical scenario, a named pipe has several writers
 - Named pipes continue to exist after communicating processes have finished
- ❑ Both UNIX and Windows systems support named pipes
 - Half-duplex in UNIX
 - Full-duplex in Windows

Outline

- ❑ Process Concept
- ❑ Process Scheduling
- ❑ Operations on Processes
- ❑ Interprocess Communication
- ❑ IPC in Shared-Memory Systems
- ❑ IPC in Message-Passing Systems
- ❑ Examples of IPC Systems
- ❑ **Communication in Client-Server Systems**
 - **Sockets**, Remote Procedure Calls

Sockets

- ❑ A **socket** is defined as an endpoint for communication

- A socket is identified by an IP address concatenated with a port number

- ❑ Sockets use a client-server architecture

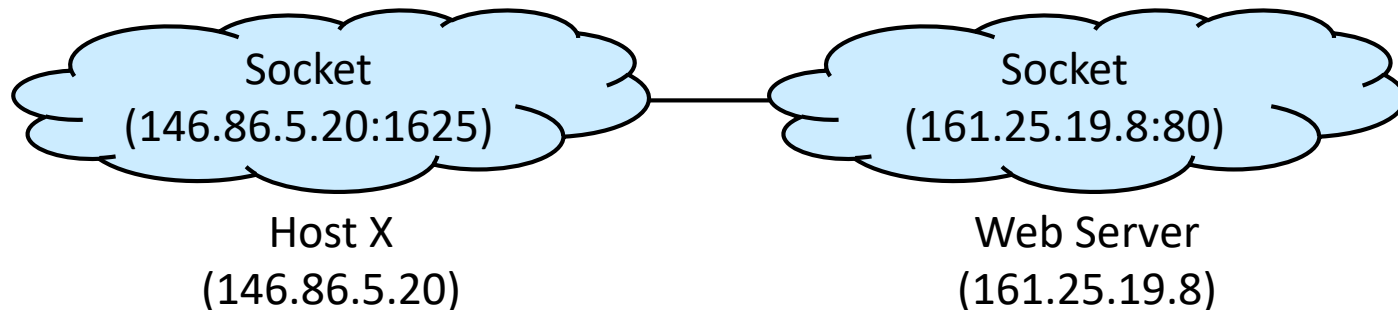
- A server waits for client requests by listening to a specified port
- Once a request is received, the server accepts a connection from the client socket to complete the connection

- ❑ All ports below 1024 are considered **well-known**

- SSH: port 22; FTP: port 21; HTTP: port 80

- ❑ Communication using a pair of sockets

- (146.86.5.20:1625) on host X and (161.25.19.8:80) on the web server



Sockets in Java

❑ Three types of sockets

- Connection-oriented (TCP) sockets
 - `Socket` class
- Connectionless (UDP) sockets
 - `DatagramSocket` class
- Multicast sockets, allowing data to be sent to multiple recipients
 - `MulticastSocket` class, a subclass of the `DatagramSocket` class

Date Server in Java

```
import java.net.*;
import java.io.*;
public class DateServer
{
    public static void main(String[] args) {
        try {
            ServerSocket sock = new ServerSocket(6013);
            /* now listen for connections */
            while (true) {
                Socket client = sock.accept();
                PrintWriter pout = new
                    PrintWriter(client.getOutputStream(), true);
                /* write the Date to the socket */
                pout.println(new java.util.Date().toString());
                /* close the socket and resume */
                /* listening for connections */
                client.close();
            }
        }
        catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```

Date Server in Java

- ❑ The server listens to the port with the **accept()** method
 - The server blocks on the **accept()** method waiting for a client to request a connection
 - When a connection request is received, **accept()** returns a socket that the server can use to communicate with the client

```
        while (true) {
            Socket client = sock.accept();
            PrintWriter pout = new
                PrintWriter(client.getOutputStream(), true);
            /* write the Date to the socket */
            pout.println(new java.util.Date().toString());
            /* close the socket and resume */
            /* listening for connections */
            client.close();
        }
    }
    catch (IOException ioe) {
        System.err.println(ioe);
    }
}
```

Date Client in Java

```
import java.net.*;
import java.io.*;
public class DateClient
{
    public static void main(String[] args) {
        try {
            /* make connection to server socket */
            Socket sock = new Socket("127.0.0.1",6013);
            InputStream in = sock.getInputStream();
            BufferedReader bin = new
                BufferedReader(new InputStreamReader(in));
            /* read the date from the socket */
            String line;
            while ( (line = bin.readLine()) != null)
                System.out.println(line);
            /* close the socket connection*/
            sock.close();
        }
        catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```

Outline

- ❑ Process Concept
- ❑ Process Scheduling
- ❑ Operations on Processes
- ❑ Interprocess Communication
- ❑ IPC in Shared-Memory Systems
- ❑ IPC in Message-Passing Systems
- ❑ Examples of IPC Systems
- ❑ **Communication in Client-Server Systems**
 - Sockets, Remote Procedure Calls

Remote Procedure Calls (1/2)

- ❑ Remote procedure call (RPC) abstracts the procedure-call mechanism between systems with network connections
 - Allow a client to invoke a procedure on a remote host as it would invoke a procedure locally
- ❑ The RPC system provides a stub on the client side
 - When the client invokes a remote procedure, the RPC system calls the appropriate stub (with parameters)
 - This stub locates the port on the server and marshals the parameters
 - The stub transmits a message to the server using message passing
 - A similar stub on the server side receives this message and invokes the procedure on the server
 - If necessary, return values are passed back to the client using the same technique
- ❑ Microsoft Interface Definition Language (MIDL) on windows

Remote Procedure Calls (2/2)

❑ Many RPC systems define a machine-independent representation of data

➤ Example: external data representation (XDR)

- On the client side, parameter marshaling involves converting the machine-dependent data into XDR
- On the server side, the XDR data are unmarshaled and converted to the machine-dependent representation

❑ Remote communication has more failure scenarios

➤ Messages can be acted on exactly once, rather than at most once

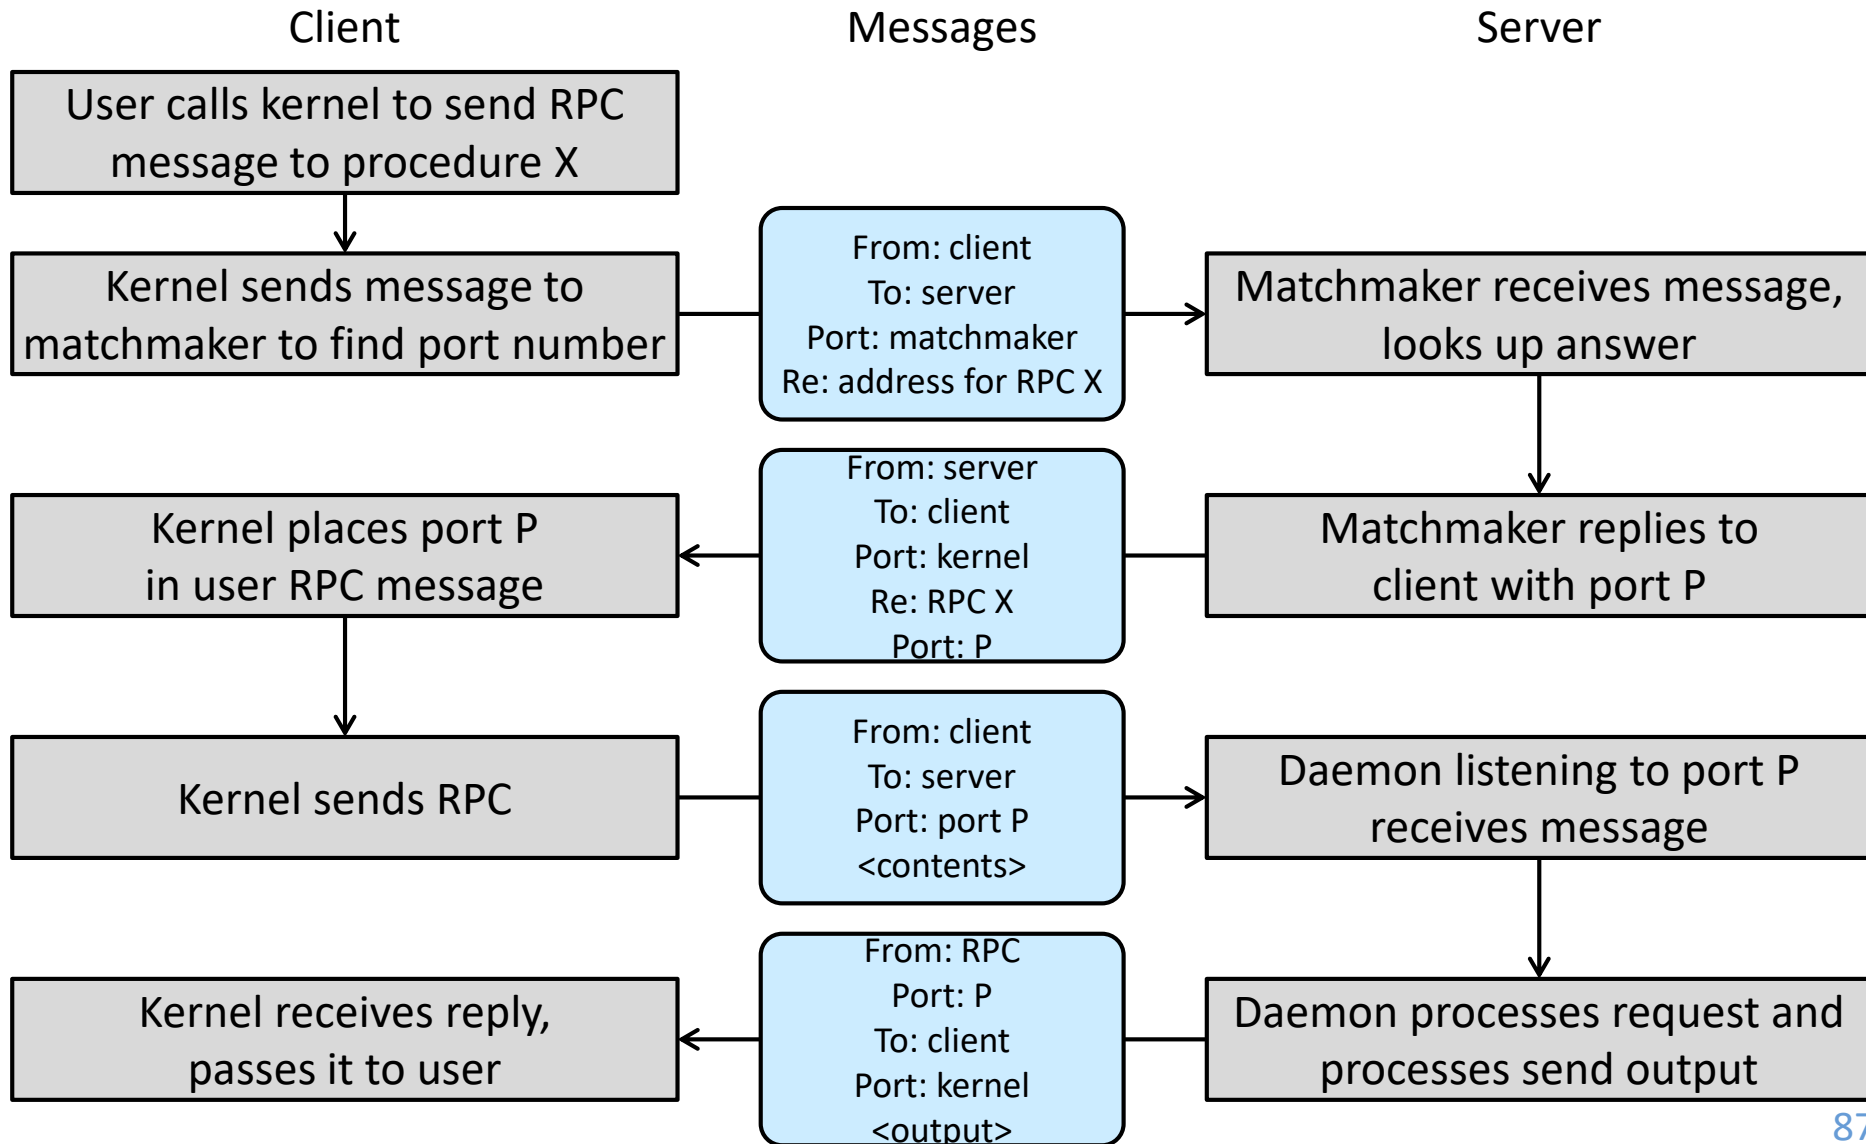
❑ Binding

➤ Predetermined

➤ Dynamic

- An operating system provides a rendezvous (also called a matchmaker) daemon on a fixed RPC port

Dynamic Binding and Execution of RPC



Objectives

- ❑ Identify the separate components of a process and illustrate how they are represented and scheduled in an operating system
- ❑ Describe how processes are created and terminated in an operating system, including developing programs using the appropriate system calls that perform these operations
- ❑ Describe and contrast interprocess communication using shared memory and message passing
- ❑ Design programs that use pipes and POSIX shared memory to perform interprocess communication
- ❑ Describe client-server communication using sockets and remote procedure calls

Q&A