

Operating Systems

[4. Threads & Concurrency]

Chung-Wei Lin

cwlin@csie.ntu.edu.tw

CSIE Department

National Taiwan University

Objectives

- ❑ Identify the basic components of a thread, and contrast threads and processes
- ❑ Describe the benefits and challenges of designing multithreaded applications
- ❑ Design multithreaded applications using the Pthreads, Java, and Windows threading APIs
- ❑ Illustrate different approaches to implicit threading, including thread pools, fork-join, and Grand Central Dispatch
- ❑ Describe how the Windows and Linux operating systems represent threads

Outline

☐ Overview

- Motivation
- Benefits

☐ Multicore Programming

☐ Multithreading Models

☐ Thread Libraries

☐ Implicit Threading

☐ Threading Issues

☐ Operating System Examples

Overview (1/2)

❑ A thread is a basic unit of CPU utilization

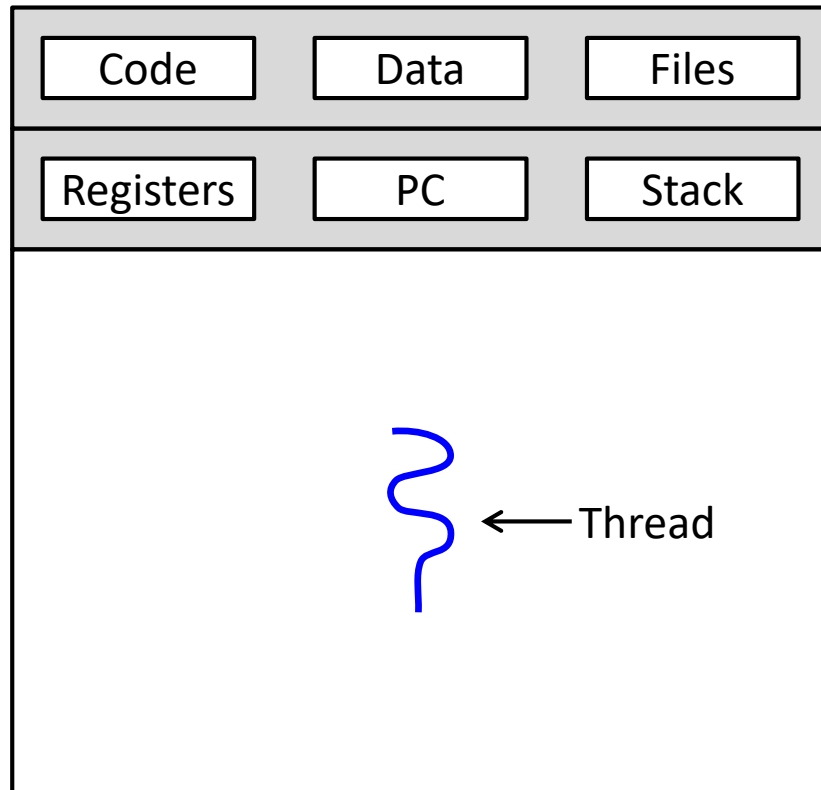
- A thread comprises
 - A thread ID
 - A program counter (PC)
 - A register set
 - A stack
- A thread shares the following items with other threads belonging to the same process
 - Code section
 - Data section
 - Other operating-system resources, such as open files and signals

❑ Thread [Wikipedia]

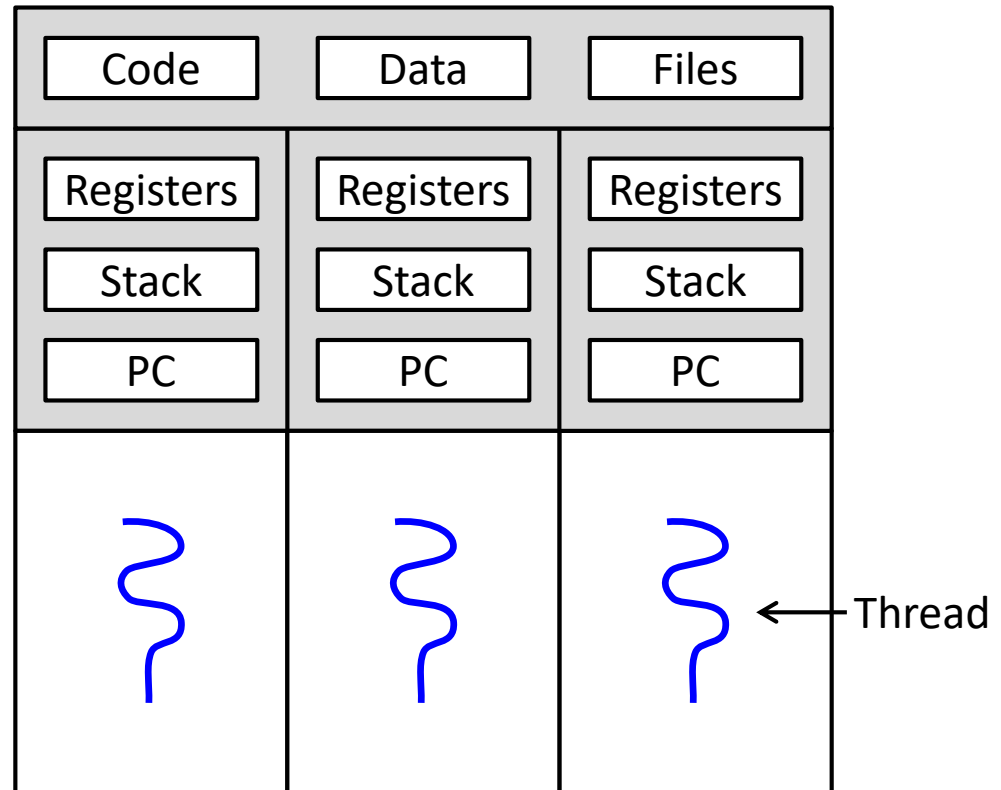
- The smallest sequence of programmed instructions that can be managed independently by a scheduler

Overview (2/2)

- ❑ A traditional process has a single thread of control
- ❑ If a process has multiple threads of control, it can perform more than one task at a time



Single-Threaded Process



Multithreaded Process

Motivation

- ❑ Most software applications that run on modern computers and mobile devices are multithreaded
 - Example: web browser
 - A thread displays images or text
 - A thread retrieves data from the network
 - Example: word processor
 - A thread displays graphics
 - A thread responds to keystrokes from the user
 - A thread performs spelling and grammar checking
 - Example: CPU-intensive tasks in parallel across the multiple cores
 - Example: web server which has several clients concurrently accessing it
- ❑ Process creation is time consuming and resource intensive
 - Using one process containing multiple threads is usually more efficient
- ❑ Most OS kernels are also typically multithreaded

Benefits

❑ Responsiveness

- Multithreading allows a program to continue running even if part of it is blocked or is performing a lengthy operation
 - Especially useful in designing user interfaces

❑ Resource sharing

- Processes can share resources only through techniques which must be explicitly arranged by the programmer

❑ Economy

- It is more economical to create and context-switch threads

❑ Scalability

- Threads may be running in parallel on different processing cores


Outline

- ❑ Overview
- ❑ **Multicore Programming**
 - Programming Challenges
 - Types of Parallelism
- ❑ Multithreading Models
- ❑ Thread Libraries
- ❑ Implicit Threading
- ❑ Threading Issues
- ❑ Operating System Examples

Multicore Programming


❑ A system with a single computing core

- Concurrency means that threads run interleavingly over time



❑ A system with multiple cores

- Concurrency means that some threads can run in parallel



❑ Concurrency and parallelism

- A concurrent system supports more than one task by allowing all the tasks to make progress
- A parallel system can perform more than one task simultaneously
- It is possible to have concurrency without parallelism


Programming Challenges

- ❑ The trend toward multicore systems continues to place pressure on system designers and application programmers
 - Identifying tasks
 - Find areas that can be divided into separate and concurrent tasks
 - Balance
 - Ensure that tasks perform equal work of equal value (or appropriate work)
 - Data splitting
 - Divide the data accessed and manipulated by tasks to run on separate cores
 - Data dependency
 - Examine the data accessed by tasks for dependencies between the tasks
 - Testing and debugging
 - Test and debug as there are many different execution paths are possible

Types of Parallelism


❑ Data parallelism

- Distribute subsets of the same data across multiple computing cores and perform the same operation on each core



❑ Task parallelism

- Distribute not data but tasks (threads) across multiple computing cores
 - Each thread is performing a unique operation



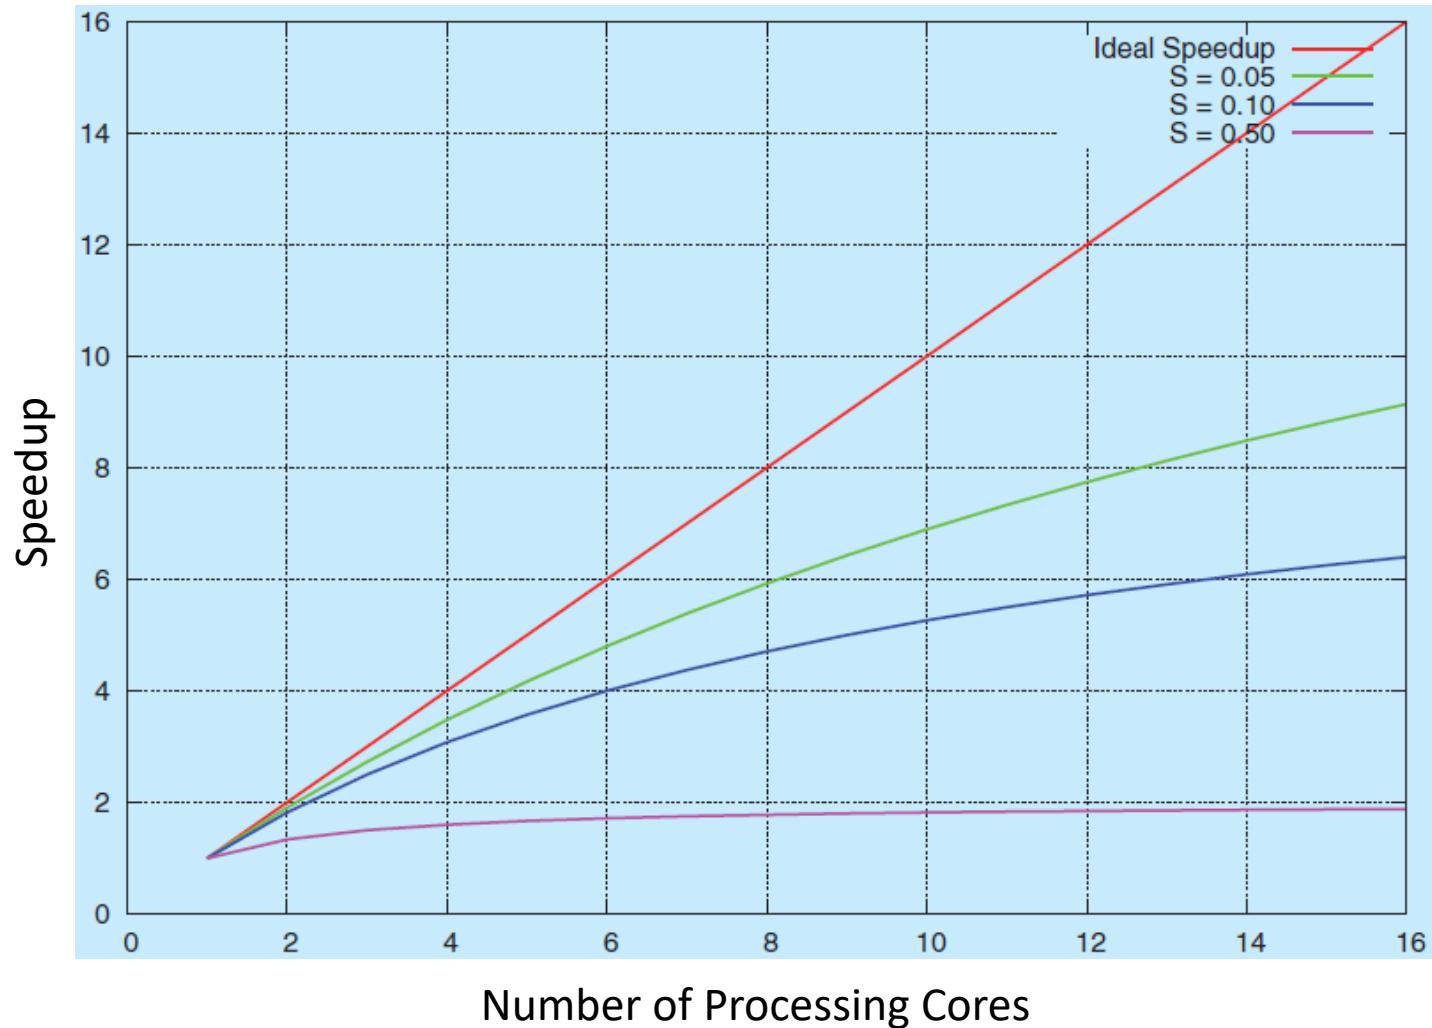
❑ They are not mutually exclusive

Amdahl's Law (1/2)

- ❑ Identify potential performance gains from adding additional computing cores to an application that has
 - Serial (nonparallel) components
 - Parallel components
- ❑ $\text{Speedup} \leq 1 / (S + \frac{1 - S}{N})$
 - S: the portion of the application that must be performed serially
 - N: the number of processing cores
- ❑ Example
 - An application that is 75 percent parallel and 25 percent serial
 - If we run this application on a system with two processing cores, we can get a speedup of 1.6 times
- ❑ Serial portion of an application has disproportionate effect on performance gained by adding additional cores

Amdahl's Law (2/2)

□ Amdahl's Law in several different scenarios



Outline

- ❑ Overview
- ❑ Multicore Programming
- ❑ **Multithreading Models**
 - Many-to-One Model
 - One-to-One Model
 - Many-to-Many Model
- ❑ Thread Libraries
- ❑ Implicit Threading
- ❑ Threading Issues
- ❑ Operating System Examples

User Threads and Kernel Threads

❑ Support for threads may be provided either

- At the user level, for user threads, or
- By the kernel, for kernel threads

❑ User threads

- Supported above the kernel and managed without kernel support
 - Primary thread libraries: POSIX Pthreads, Windows threads, Java threads

❑ Kernel threads


- Supported and managed directly by the operating system
 - Virtually all contemporary operating systems support kernel threads:
Windows, Linux, macOS, iOS, Android

❑ Ultimately, a relationship must exist between user threads and kernel threads

Many-to-One Model

❑ Map many user-level threads to one kernel thread


- The entire process will block if a thread makes a blocking system call
- Only one thread can access the kernel at a time
 - Multiple threads are unable to run in parallel on multicore systems
 - Very few systems continue to use the model
- Thread management is done by the thread library in user space
 - Therefore, it is efficient



One-to-One Model


❑ Map each user thread to a kernel thread

- Other threads can run when a thread makes a blocking system call
- Multiple threads can run in parallel on multiprocessors
- Creating a user thread requires creating the corresponding kernel thread
 - A large number of kernel threads may burden the performance of a system
- Linux, along with the family of Windows, implement this



Many-to-Many Model (1/2)

- ❑ Multiplex many user-level threads to a smaller or equal number of kernel threads
 - The number of kernel threads may be specific to either a particular application or a particular machine
 - Developers can create as many user threads as necessary, and the corresponding kernel threads can run in parallel on a multiprocessor
 - Other threads can run when a thread makes a blocking system call




Many-to-Many Model (2/2)

❑ Two-level model, one variation on the many-to-many model

- Multiplex many user level threads to a smaller or equal number of kernel threads
- Also allow a user-level thread to be bound to a kernel thread

❑ The many-to-many model is the most flexible here

- It is difficult to implement
- Limiting the number of kernel threads has become less important
 - There is an increasing number of processing cores appearing on most systems



Threading Models [Wikipedia]

❑ Thread [Wikipedia]

- The smallest sequence of programmed instructions that can be managed independently by a scheduler

❑ Many-to-one model (user-level threading)

- All application-level threads map to one kernel-level scheduled entity

❑ One-to-one model (kernel-level threading)

- Threads created by the user in a 1:1 correspondence with schedulable entities in the kernel

❑ Many-to-many model (hybrid threading)

- Map M application threads onto N kernel entities (or virtual processors)

Outline

- ❑ Overview
- ❑ Multicore Programming
- ❑ Multithreading Models
- ❑ **Thread Libraries**
 - Pthreads
 - Windows Threads
 - Java Threads
- ❑ Implicit Threading
- ❑ Threading Issues
- ❑ Operating System Examples

Thread Libraries (1/2)

- ❑ A **thread library** provides the programmer with an API for creating and managing threads
- ❑ Approach 1
 - Provide a library entirely in user space with no kernel support
 - Code and data structures for the library exist in user space
 - Invoking a function in the library results in a local function call in user space, not a system call
- ❑ Approach 2
 - Implement a kernel-level library supported directly by the operating system
 - Code and data structures for the library exist in kernel space
 - Invoking a function in the API for the library typically results in a system call to the kernel

Thread Libraries (2/2)

❑ Running example: summation from 1 to N

- Input N = 5
- Output sum = 15

❑ Asynchronous threading

- Once the parent creates a child thread, the parent resumes its execution
 - The parent and child execute concurrently and independently of one another

❑ Synchronous threading

- The parent thread creates one or more children and then must wait for all of its children to terminate before it resumes
 - The threads created by the parent perform work concurrently, but the parent cannot continue until this work has been completed

❑ All of the following examples use synchronous threading

Outline

- ❑ Overview
- ❑ Multicore Programming
- ❑ Multithreading Models
- ❑ **Thread Libraries**
 - **Pthreads**
 - Windows Threads
 - Java Threads
- ❑ Implicit Threading
- ❑ Threading Issues
- ❑ Operating System Examples

Pthreads

- ❑ A POSIX standard (IEEE 1003.1c) defining an API for thread creation and synchronization
 - This is a **specification** for thread behavior, not an implementation
 - Numerous systems implement the Pthreads specification
 - Most are UNIX-type systems, including Linux and macOS
 - Threads may be provided as either a user-level or a kernel-level library

Pthreads API

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
.....
int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* threads call this function */
.....
int main(int argc, char *argv[]) {
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */
    /* set the default attributes of the thread */
    pthread_attr_init(&attr);
    /* create the thread */
    pthread_create(&tid, &attr, runner, argv[1]);
    /* wait for the thread to exit */
    pthread_join(tid, NULL);
    printf("sum = %d\n", sum);
}
.....
/* The thread will execute in this function */
void *runner(void *param) {
    int i, upper = atoi(param);
    sum = 0;
    for (i = 1; i <= upper; i++)
        sum += i;
    pthread_exit(0);
}
```

Pthreads API

```
#include <pthread.h>
```

❑ `pthread_attr_init()` sets the attributes

- Each thread has a set of attributes, including stack size and scheduling information
- We use the default attributes provided

```
/* set the default attributes of the thread */
pthread_attr_init(&attr);
/* create the thread */
pthread_create(&tid, &attr, runner, argv[1]);
/* wait for the thread to exit */
pthread_join(tid, NULL);
printf("sum = %d\n", sum);
}
.....
/* The thread will execute in this function */
void *runner(void *param) {
    int i, upper = atoi(param);
    sum = 0;
    for (i = 1; i <= upper; i++)
        sum += i;
    pthread_exit(0);
}
```

Pthreads API

```
#include <pthread.h>
```

❑ **pthread_create()** creates a separate thread with

- The thread identifier
- The attributes for the thread
- The name of the function where the new thread will begin execution
- The integer parameter

```
/* create the thread */
pthread_create(&tid, &attr, runner, argv[1]);
/* wait for the thread to exit */
pthread_join(tid, NULL);
printf("sum = %d\n", sum);
}
.....
/* The thread will execute in this function */
void *runner(void *param) {
    int i, upper = atoi(param);
    sum = 0;
    for (i = 1; i <= upper; i++)
        sum += i;
    pthread_exit(0);
}
```

Pthreads API

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
.....
int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* threads call this function */
.....
int main(int argc, char *argv[]) {
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */
```

❑ The parent thread waits for the summation (child) thread to terminate by calling the **pthread_join()** function

```
    /* wait for the thread to exit */
    pthread_join(tid, NULL);
    printf("sum = %d\n", sum);
}
.....
/* The thread will execute in this function */
void *runner(void *param) {
    int i, upper = atoi(param);
    sum = 0;
    for (i = 1; i <= upper; i++)
        sum += i;
    pthread_exit(0);
}
```

Pthreads API

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
.....
int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* threads call this function */
.....
int main(int argc, char *argv[]) {
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */
```

❑ The summation thread terminates when it calls the function **pthread_exit()**

```
    /* wait for the thread to exit */
    pthread_join(tid, NULL);
    printf("sum = %d\n", sum);
}
.....
/* The thread will execute in this function */
void *runner(void *param) {
    int i, upper = atoi(param);
    sum = 0;
    for (i = 1; i <= upper; i++)
        sum += i;
    pthread_exit(0);
}
```

Comparison with Process Creation

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
int main() {
    pid_t pid;
    /* fork a child process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }
    return 0;
}
```

Comparison with Function Call

```
#include <stdio.h>
#include <stdlib.h>

int sum;
void runner(void *param);

int main(int argc, char *argv[]) {
    runner(argv[1]);
    printf("sum = %d\n", sum);
}

void runner(void *param) {
    int i, upper = atoi(param);
    sum = 0;
    for (i = 1; i <= upper; i++)
        sum += i;
}
```


Waiting on Several Threads

❑ Pthread code for joining 10 threads

```
#define NUM_THREADS 10

/* an array of threads to be joined upon */
pthread_t workers[NUM_THREADS];

for (int i = 0; i < NUM_THREADS; i++)
    pthread_join(workers[i], NULL);
```

Outline

- ❑ Overview
- ❑ Multicore Programming
- ❑ Multithreading Models
- ❑ **Thread Libraries**
 - Pthreads
 - **Windows Threads**
 - Java Threads
- ❑ Implicit Threading
- ❑ Threading Issues
- ❑ Operating System Examples

Windows Threads

```
#include <windows.h>
#include <stdio.h>
DWORD Sum; /* data is shared by the thread(s) */
/* The thread will execute in this function */
DWORD WINAPI Summation(LPVOID Param) {
    DWORD Upper = *(DWORD*)Param;
    for (DWORD i = 1; i <= Upper; i++)
        Sum += i;
    return 0;
}
int main(int argc, char *argv[]) {
    DWORD ThreadId;
    HANDLE ThreadHandle;
    int Param;
    Param = atoi(argv[1]);
    /* create the thread */
    ThreadHandle =
        CreateThread(NULL, 0, Summation, &Param, 0, &ThreadId);
    /* now wait for the thread to finish */
    WaitForSingleObject(ThreadHandle, INFINITE);
    /* close the thread handle */
    CloseHandle(ThreadHandle);
    printf("sum = %d\n", Sum);
}
```

Waiting on Several Threads

❑ **WaitForMultipleObjects(N, THandles, TRUE, INFINITE) ;**

- The number of objects to wait for
- A pointer to the array of objects
- A flag indicating whether all objects have been signaled
- A timeout duration (or **INFINITE**)

Outline

- ❑ Overview
- ❑ Multicore Programming
- ❑ Multithreading Models
- ❑ **Thread Libraries**
 - Pthreads
 - Windows Threads
 - **Java Threads**
- ❑ Implicit Threading
- ❑ Threading Issues
- ❑ Operating System Examples

Java Threads (1/3)

- ❑ Java threads are available on any system that provides a Java virtual machine (JVM) including Windows, Linux, and macOS
 - The Java thread API is available for Android applications as well
 - The Java thread API is generally implemented using a thread library available on the host system
- ❑ Two techniques for explicitly creating threads
 - Create a new class derived from the **Thread** class and override its **run()** method
 - The (default) **run()** method is called if the thread was constructed using a separate **Runnable** object; otherwise, this method does nothing and returns
 - The **run()** method can be called using the **start()** method
 - Define a class that implements the **Runnable** interface
 - **Runnable** is an interface used to execute code on a concurrent thread
 - More common

Java Threads (2/3)

❑ Define a class that implements the **Runnable** interface

```
class Task implements Runnable {  
    public void run() {  
        System.out.println("I am a thread.");  
    }  
}
```

- The code in the **run()** method is what executes in a separate thread

❑ Create a thread

```
Thread worker = new Thread(new Task());  
worker.start();
```

- Invoking the **start()** method does two things
 - Allocate memory and initialize a new thread in the JVM
 - Call the **run()** method, making the thread eligible to run by the JVM
- We call the **start()** method, not the **run()** method

Java Threads (3/3)

❑ Wait for a thread

```
try {  
    worker.join();  
}  
catch (InterruptedException ie) { }
```


Java **Executor** Framework

- ❑ Beginning with Version 1.5 and its API, Java provide greater control over thread creation and communication
 - Available in the `java.util.concurrent` package
- ❑ Classes implementing the **Executor** interface must define the **execute ()** method
 - The **execute ()** method is passed a **Runnable** object
 - For Java developers, this means using the **Executor** rather than creating a separate **Thread** object and invoking its **start ()** method

Outline

- ❑ Overview
- ❑ Multicore Programming
- ❑ Multithreading Models
- ❑ Thread Libraries
- ❑ **Implicit Threading**
 - Thread Pools
 - Fork Join
 - OpenMP
 - Grand Central Dispatch
 - Intel Thread Building Blocks
- ❑ Threading Issues
- ❑ Operating System Examples

Implicit Threading

- ❑ With the continued growth of multicore processing, designing applications containing many threads is difficult
- ❑ **Implicit threading**
 - Transfer the creation and management of threading from application developers to compilers and run-time libraries
- ❑ Alternative approaches to designing applications that can take advantage of multicore processors through implicit threading
 - These strategies generally require application developers to identify **tasks** (not threads) that can run in parallel
 - A task is usually written as a function
 - The run-time library maps it to a separate thread, typically using the many-to-many model

Outline

- ❑ Overview
- ❑ Multicore Programming
- ❑ Multithreading Models
- ❑ Thread Libraries
- ❑ **Implicit Threading**
 - **Thread Pools**
 - Fork Join
 - OpenMP
 - Grand Central Dispatch
 - Intel Thread Building Blocks
- ❑ Threading Issues
- ❑ Operating System Examples

Thread Pools

❑ Create a number of threads and place them into a pool

- They sit and wait for work

❑ Example: a server

- When a server receives a request, it submits the request to the thread pool and resumes waiting for additional requests
 - The server does not create a thread
- If there is an available thread in the pool, it is awakened, and the request is serviced immediately
- If the pool contains no available thread, the task is queued until one becomes free
- Once a thread completes its service, it returns to the pool and awaits more work

❑ Demo [Prof. Shih]

Thread Pools Advantages

- ❑ Servicing a request with an existing thread is often faster than waiting to create a thread
- ❑ A thread pool limits the number of threads that exist at any one point
 - Important on systems that cannot support a large number of concurrent threads
- ❑ A thread pool allows us to use different strategies for running the task
 - Separate the task to be performed from the mechanics of creating the task
 - Example: the task could be scheduled to execute after a time delay or to execute periodically

Windows Thread Pool API

- ❑ A example function that is to run as a separate thread

```
DWORD WINAPI PoolFunction(PVOID Param) {  
    /* this function runs as a separate thread */  
}
```

- ❑ An example of invoking a function

```
QueueUserWorkItem(&PoolFunction, NULL, 0);
```

- A pointer to the function that is to run as a separate thread
- The parameter passed to the function
- The flags indicating how the thread pool is to create and manage execution of the thread

Java Thread Pools

- ❑ Single thread executor creates a pool of size 1
 - `newSingleThreadExecutor()`
- ❑ Fixed thread executor creates a thread pool with a specified number of threads
 - `newFixedThreadPool(int size)`
- ❑ Cached thread executor creates an unbounded thread pool, reusing threads in many instances
 - `newCachedThreadPool()`

Outline

- ❑ Overview
- ❑ Multicore Programming
- ❑ Multithreading Models
- ❑ Thread Libraries
- ❑ **Implicit Threading**
 - Thread Pools
 - **Fork Join**
 - OpenMP
 - Grand Central Dispatch
 - Intel Thread Building Blocks
- ❑ Threading Issues
- ❑ Operating System Examples

Fork Join


❑ The fork-join model

➤ Covered above

- The main parent thread creates (forks) one or more child threads and then waits for the children to terminate and join with it
- This synchronous model is often characterized as explicit thread creation

➤ It is also an excellent candidate for implicit threading

- Threads are not constructed directly during the fork stage
- Parallel tasks are designated, instead



➤ A library manages the number of threads that are created and is also responsible for assigning tasks to threads

- In some ways, this is a synchronous version of thread pools in which a library determines the actual number of threads to create

Fork Join in Java

- ❑ A fork-join library in Version 1.7 of the API that is designed to be used with recursive divide-and-conquer algorithms

```
Task(problem)
```

```
    if problem is small enough
```

```
        solve the problem directly
```

```
    else
```


```
        subtask1 = fork(new Task(subset of problem))
```

```
        subtask2 = fork(new Task(subset of problem))
```

```
        result1 = join(subtask1)
```

```
        result2 = join(subtask2)
```

```
    return combined results
```



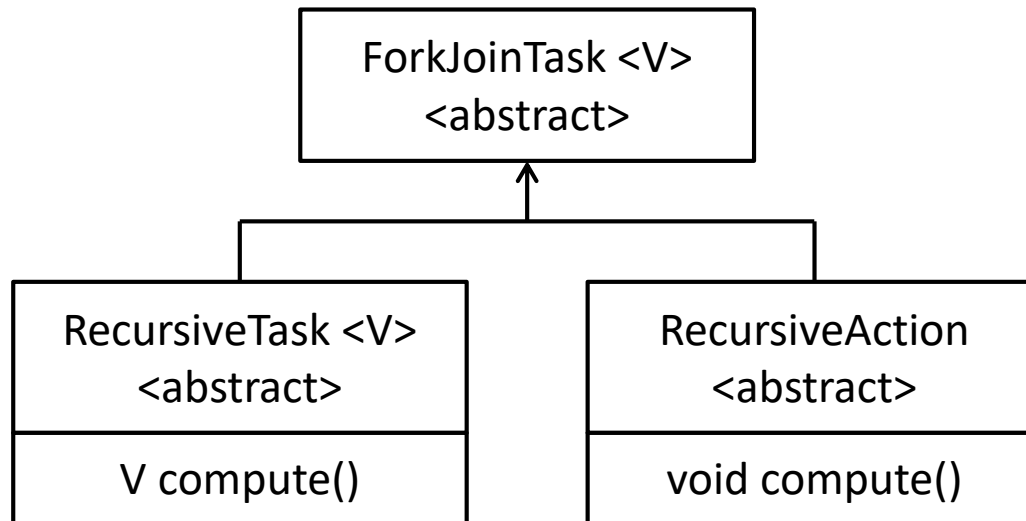
Fork Join in Java: Example (1/2)

```
protected Integer compute() {
    if (end - begin < THRESHOLD) {
        int sum = 0;
        for (int i = begin; i <= end; i++)
            sum += array[i];
        return sum;
    }
    else {
        int mid = (begin + end) / 2;
        SumTask leftTask = new SumTask(begin, mid, array);
        SumTask rightTask = new SumTask(mid + 1, end, array);
        leftTask.fork();
        rightTask.fork();
        return rightTask.join() + leftTask.join();
    }
}
```

Fork Join in Java: Example (2/2)

❑ **ForkJoinTask** is an abstract base class

- **RecursiveTask** and **RecursiveAction** classes extend it
- **RecursiveTask** returns a result
 - Via the return value from the **compute()** method
- **RecursiveAction** does not return a result



Outline

- ❑ Overview
- ❑ Multicore Programming
- ❑ Multithreading Models
- ❑ Thread Libraries
- ❑ **Implicit Threading**
 - Thread Pools
 - Fork Join
 - **OpenMP**
 - Grand Central Dispatch
 - Intel Thread Building Blocks
- ❑ Threading Issues
- ❑ Operating System Examples

OpenMP

- ❑ A set of compiler directives and an API for C, C++, or FORTRAN

- Support parallel programming in shared memory environments

- ❑ Developers insert compiler directives at parallel regions

```
#pragma omp parallel
{
    printf("I am a parallel region.");
}

#pragma omp parallel for
for (i = 0; i < N; i++) {
    c[i] = a[i] + b[i];
}
```

- Instruct the OpenMP run-time library to execute the region in parallel
 - OpenMP creates as many threads as there are processing cores in the system

- ❑ Demo [Prof. Shih]

Outline

- ❑ Overview
- ❑ Multicore Programming
- ❑ Multithreading Models
- ❑ Thread Libraries
- ❑ **Implicit Threading**
 - Thread Pools
 - Fork Join
 - OpenMP
 - **Grand Central Dispatch**
 - Intel Thread Building Blocks
- ❑ Threading Issues
- ❑ Operating System Examples

Grand Central Dispatch (GCD) (1/2)

❑ A technology developed by Apple for its macOS and iOS

- Include a run-time library, an API, and language extensions
- Allow developers to identify sections of code (tasks) to run in parallel
- Manage most of the details of threading (like OpenMP)

❑ GCD schedules tasks for run-time execution by placing them on a dispatch queue

- Serial (private) dispatch queue
 - Each process has its own serial queue
 - Once a task has been removed from the queue, it must complete execution before another task is removed
- Concurrent (global) dispatch queue
 - Several tasks may be removed at a time
 - Several system-wide concurrent queues are divided into four primary quality-of-service classes: user-interactive, user-initiated, utility, and background

Grand Central Dispatch (GCD) (2/2)

❑ Two different ways to express tasks submitted to dispatch queues

- A language extension, block, for C, C++, and Objective-C languages

```
^ { printf("I am a block"); }
```

- A closure, similar to a block, for the Swift programming language

```
dispatch_async(queue, { print("A closure.") })
```

Outline

- ❑ Overview
- ❑ Multicore Programming
- ❑ Multithreading Models
- ❑ Thread Libraries
- ❑ **Implicit Threading**
 - Thread Pools
 - Fork Join
 - OpenMP
 - Grand Central Dispatch
 - **Intel Thread Building Blocks**
- ❑ Threading Issues
- ❑ Operating System Examples

Intel Threading Building Blocks (TBB)

- ❑ A template library supporting parallel applications in C++
- ❑ A serial for loop

```
for (int i = 0; i < n; i++) {  
    apply(v[i]);  
}
```

- ❑ The for loop using the TBB **parallel_for** template

```
parallel_for (size_t(0), n, [=](size_t i) {apply(v[i]);});
```

Outline

- ❑ Overview
- ❑ Multicore Programming
- ❑ Multithreading Models
- ❑ Thread Libraries
- ❑ Implicit Threading
- ❑ **Threading Issues**
 - **The `fork()` and `exec()` System Calls**
 - Signal Handling
 - Thread Cancellation
 - Thread-Local Storage
 - Scheduler Activations
- ❑ Operating System Examples

Recap: Process Creation

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
int main() {
    pid_t pid;
    /* fork a child process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }
    return 0;
}
```

fork () and **exec ()** System Calls

- ❑ The semantics of the **fork ()** and **exec ()** system calls change in a multithreaded program
 - Some UNIX systems have chosen to have two versions of **fork ()**
 - One duplicates all threads
 - One duplicates only the thread that invoked the **fork()** system call
 - The **exec ()** system call typically works in the same way
 - Replace the entire process including all threads
- ❑ If **exec ()** is called immediately after forking, then duplicating all threads is unnecessary

Outline

- ❑ Overview
- ❑ Multicore Programming
- ❑ Multithreading Models
- ❑ Thread Libraries
- ❑ Implicit Threading
- ❑ Threading Issues
 - The **fork()** and **exec()** System Calls
 - Signal Handling
 - Thread Cancellation
 - Thread-Local Storage
 - Scheduler Activations
- ❑ Operating System Examples

Signals

❑ A signal is used in UNIX systems to notify a process

- All signals follow the same pattern
 - A signal is generated by the occurrence of a particular event
 - The signal is delivered to a process
 - Once delivered, the signal must be handled?!

❑ Synchronous signal

- Synchronous signals are delivered to the same process that performed the operation that caused the signal
- Examples: illegal memory access and division by 0

❑ Asynchronous signal

- The signal is generated by an event external to a running process
- Examples: terminating a process with Ctrl+C and a timer expired

Signal Handling

❑ Two possible handlers

- Every signal has a **default signal handler** that the kernel runs when handling that signal
- This default action can be overridden by a **user-defined signal handler**

❑ Signals are handled in different ways

- Examples: terminating the program, ignoring the signal

❑ For single-threaded programs

- Signals are always delivered to a process

❑ For multithreaded programs

- Deliver the signal to the thread to which the signal applies
- Deliver the signal to every thread in the process
- Deliver the signal to certain threads in the process
- Assign a specific thread to receive all signals for the process

Outline

- ❑ Overview
- ❑ Multicore Programming
- ❑ Multithreading Models
- ❑ Thread Libraries
- ❑ Implicit Threading
- ❑ Threading Issues
 - The **fork()** and **exec()** System Calls
 - Signal Handling
 - Thread Cancellation
 - Thread-Local Storage
 - Scheduler Activations
- ❑ Operating System Examples

Thread Cancellation

- ❑ **Thread cancellation** involves terminating a thread before it has completed
 - A thread that is to be canceled is often referred to as the **target thread**
- ❑ **Two scenarios of cancellation of a target thread**
 - Asynchronous cancellation
 - One thread immediately terminates the target thread
 - Deferred cancellation
 - The target thread periodically checks whether it should terminate, allowing it an opportunity to terminate itself in an orderly fashion
 - Canceling a thread asynchronously may not free a necessary system-wide resource

Pthread Cancellation (1/2)

❑ Thread cancellation is initiated using `pthread_cancel()`

- Only a request to cancel the target thread
- Actual cancellation depending on how the target thread is set up to
 - Three cancellation modes: disabled, deferred, asynchronous

❑ When the target thread is finally canceled, the call to `pthread_join()` in the canceling thread returns

```
pthread_t tid;  
/* create the thread */  
pthread_create(&tid, 0, worker, NULL);  
/* cancel the thread */  
pthread_cancel(tid);  
/* wait for the thread to terminate */  
pthread_join(tid, NULL);
```

Pthread Cancellation (2/2)

❑ Default cancellation mode: deferred cancellation

- However, cancellation occurs only when a thread reaches a cancellation point which can be set by `pthread_testcancel()`

```
while (1) {  
    /* do some work for awhile */  
  
    ...  
  
    /* is there a cancellation request? */  
    pthread_testcancel();  
}
```

❑ A cleanup handler can be invoked if a thread is canceled

- Release any resource that the thread may have acquired

❑ On Linux systems, thread cancellation using the Pthreads API is handled through signals

Thread Cancellation in Java

❑ A policy similar to deferred cancellation in Pthreads

- The `interrupt()` method set the interruption status of the target thread to true

```
Thread worker;
```

```
. . .
```

```
/* set the interruption status of the thread */  
worker.interrupt()
```

- A thread can check its interruption status by invoking the `isInterrupted()` method

```
while(!Thread.currentThread().isInterrupted()) {  
    . . .  
}
```

Outline

- ❑ Overview
- ❑ Multicore Programming
- ❑ Multithreading Models
- ❑ Thread Libraries
- ❑ Implicit Threading
- ❑ Threading Issues
 - The **fork()** and **exec()** System Calls
 - Signal Handling
 - Thread Cancellation
 - Thread-Local Storage
 - Scheduler Activations
- ❑ Operating System Examples

Thread-Local Storage

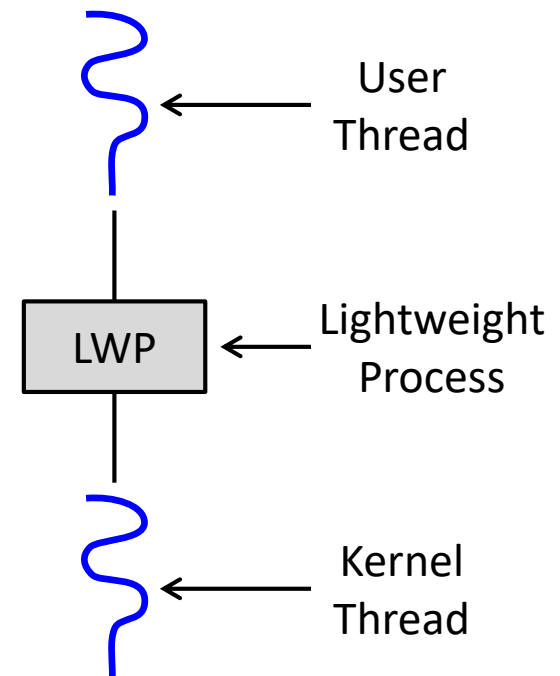
- ❑ Thread-local storage (TLS) allows each thread to have its own copy of certain data
 - Different from local variables
 - Local variables are visible only during a single function invocation
 - TLS data are visible across function invocations
 - Useful the developer has no control over the thread creation process
 - Example: when using an implicit technique such as a thread pool
 - Similar to **static** data
 - A static variable inside a function keeps its value between invocations
 - However, TLS is unique to each thread

Outline

- ❑ Overview
- ❑ Multicore Programming
- ❑ Multithreading Models
- ❑ Thread Libraries
- ❑ Implicit Threading
- ❑ Threading Issues
 - The **fork()** and **exec()** System Calls
 - Signal Handling
 - Thread Cancellation
 - Thread-Local Storage
 - Scheduler Activations
- ❑ Operating System Examples

Lightweight Process (LWP)

- ❑ Many-to-many and two-level models require communication to dynamically adjust the number of kernel threads
- ❑ An intermediate data structure, an LWP, is placed between the user and kernel threads by many systems
 - An LWP appears to be a virtual processor to the user-thread library
 - An application can schedule a user thread to run on it
 - Each LWP is attached to a kernel thread
 - An operating system schedules kernel threads to run on physical processors



Scheduler Activations

- ❑ One scheme for communication between the user-thread library and the kernel
 - The kernel provides an application with a set of LWPs
 - The application can schedule user threads onto an available LWP
- ❑ The kernel must inform an application about certain events
 - This procedure is known as an upcall
 - Upcalls are handled by the thread library with an upcall handler
 - Upcall handlers must run on a virtual processor

Outline

- ❑ Overview
- ❑ Multicore Programming
- ❑ Multithreading Models
- ❑ Thread Libraries
- ❑ Implicit Threading
- ❑ Threading Issues
- ❑ **Operating System Examples**
 - Windows Threads
 - Linux Threads

Windows Threads (1/2)

- ❑ A Windows application runs as a separate process
- ❑ Each process may contain one or more threads
 - Windows API for creating threads has been covered
 - Windows uses the one-to-one mapping model
- ❑ The general components of a thread
 - A thread ID uniquely identifying the thread
 - A register set representing the status of the processor
 - A program counter
 - User and kernel stacks, employed when the thread is running in user and kernel modes, respectively
 - A private storage area used by run-time and dynamic link libraries
- ❑ The register set, stacks, and private storage area are known as the context of the thread

Windows Threads (2/2)

❑ The primary data structures of a thread

- **ETHREAD**: executive thread block (in kernel space)
 - A pointer to the process to which the thread belongs
 - The address of the routine in which the thread starts control
 - A pointer to the corresponding **KTHREAD**
- **KTHREAD**: kernel thread block (in kernel space)
 - Scheduling and synchronization information for the thread
 - The kernel stack used when the thread is running in kernel mode
 - A pointer to the **TEB**
- **TEB**: thread environment block (in user space)
 - The thread identifier
 - A user-mode stack
 - An array for thread-local storage

Linux Threads

- ❑ Linux provides the ability to create threads using **clone()**
 - In fact, Linux uses the term **task**, rather than process or thread, when referring to a flow of control within a program
- ❑ **clone()** is passed a set of flags that determine how much sharing is to take place between the parent and child tasks
 - **CLONE_FS**: file-system information is shared
 - **CLONE_VM**: the same memory space is shared
 - **CLONE_SIGHAND**: signal handlers are shared
 - **CLONE_FILES**: the set of open files is shared
 - If none of these flags is set, no sharing takes place, resulting in functionality similar to that provided by **fork()**
- ❑ A unique kernel data structure (**struct task_struct**) exists for each task in the system

Objectives

- ❑ Identify the basic components of a thread, and contrast threads and processes
- ❑ Describe the benefits and challenges of designing multithreaded applications
- ❑ Design multithreaded applications using the Pthreads, Java, and Windows threading APIs
- ❑ Illustrate different approaches to implicit threading, including thread pools, fork-join, and Grand Central Dispatch
- ❑ Describe how the Windows and Linux operating systems represent threads

Q&A