

Operating Systems

[14A. File-System Implementation]

Chung-Wei Lin

cwlin@csie.ntu.edu.tw

CSIE Department

National Taiwan University

Objectives

- ❑ Describe the details of implementing local file systems and directory structures
- ❑ Discuss block allocation and free-block algorithms and trade-offs
- ❑ Explore file system efficiency and performance issues
- ❑ Look at recovery from file system failures
- ❑ Describe the WAFL file system as a concrete example

Outline

- ❑ **File-System Structure**
- ❑ File-System Operations
- ❑ Directory Implementation
- ❑ Allocation Methods
- ❑ Free-Space Management
- ❑ Efficiency and Performance
- ❑ Recovery
- ❑ Example: WAFL File System

File-System Structure

- ❑ To improve I/O efficiency, I/O transfers between memory and mass storage are performed in units of **blocks**
 - Each block on a hard disk drive has one or more sectors
 - Depending on the disk drive, sector size is usually 512 or 4,096 bytes
- ❑ **File systems** provide access to the storage device by allowing data to be stored, located, and retrieved easily
 - Define a file and its attributes, the operations allowed on a file, and the directory structure for organizing files
 - How the file system should look to the user
 - Create algorithms and data structures to map the logical file system onto the physical secondary-storage devices

Layered File Systems (1/3)

❑ Devices

❑ I/O control

- Device drivers and interrupt handlers to transfer information between the main memory and the disk system
- A device driver, as a translator, usually writes specific bit patterns to special locations in the I/O controller's memory
 - Its input consists of high-level commands, such as "retrieve block 123"
 - Its output consists of low-level, hardware-specific instructions that are used by the hardware controller

Application Programs
Logical File System
File-Organization Module
Basic File System
I/O control
Devices

Layered File Systems (2/3)

Application Programs
Logical File System
File-Organization Module
Basic File System
I/O control
Devices

❑ Basic file system (Linux block I/O subsystem)

- Need only to issue generic commands to the appropriate device driver to read and write blocks on the storage device
 - Issue commands to the drive based on logical block addresses
- Also manage the memory buffers and caches that hold various filesystem, directory, and data blocks
 - A block in the buffer is allocated before the transfer of a mass storage block can occur
 - When the buffer is full, the buffer manager must find more buffer memory or free up buffer space to allow a requested I/O to complete
 - Caches are used to hold frequently used file-system metadata to improve performance

Layered File Systems (3/3)

Application Programs
Logical File System
File-Organization Module
Basic File System
I/O control
Devices

❑ File-organization module

- Know about files and their logical blocks
 - Each file's logical blocks are numbered from 0 (or 1) through N
- Also include the free-space manager
 - Track unallocated blocks
 - Provide these blocks to the file-organization module when requested

❑ Logical file system

- Manage metadata which includes all of the file-system structure except the actual data (or contents of the files)
- Manage the directory structure
- Maintain file structure via **file-control blocks (FCBs)** (**inodes** in UNIX)
 - Information about the file, including ownership, permissions, and location of the file contents

❑ Application programs

Typical File-Control Block

- ❑ File permissions
- ❑ File dates (create, access, write)
- ❑ File owner, group, access-control list (ACL)
- ❑ File size
- ❑ File data blocks or pointers to file data blocks

Advantage and Disadvantage

❑ Advantage of layered file systems

- Duplication of code is minimized
 - The I/O control and sometimes the basic file-system code can be used by multiple file systems
 - Each file system can then have its own logical file-system and file-organization modules

❑ Disadvantage of layered file systems

- More operating-system overhead, which may result in decreased performance

❑ A major challenge in designing new systems

- The use of layering, including the decision about how many layers to use and what each layer should do

File Systems in Use

- ❑ Most CD-ROMs are written in the ISO 9660 format
- ❑ UNIX uses the UNIX file system (UFS), which is based on the Berkeley Fast File System (FFS)
- ❑ Windows supports disk file-system formats of FAT, FAT32, and NTFS, as well as CD-ROM and DVD file-system formats
- ❑ The standard Linux file system is known as the **extended file system**, with the most common versions being ext3 and ext4
 - Although Linux supports over 130 different file systems
- ❑ More are coming
 - ZFS, GoogleFS, Oracle ASM, FUSE

Outline

- ❑ File-System Structure
- ❑ **File-System Operations**
- ❑ Directory Implementation
- ❑ Allocation Methods
- ❑ Free-Space Management
- ❑ Efficiency and Performance
- ❑ Recovery
- ❑ Example: WAFL File System

Structures on File Systems (1/2)

- ❑ A **boot control block** (per volume) contains information needed to boot an operating system from that volume
 - If the disk does not contain an operating system, this block can be empty
 - It is typically the first block of a volume
 - **Boot block** in UFS and **partition boot sector** in NTFS
- ❑ A **volume control block** (per volume) contains volume details
 - The number of blocks in the volume, the size of the blocks, a free-block count and free-block pointers, and a free-FCB count and FCB pointers
 - **Superblock** in UFS and stored **master file table** in NTFS

Structures on File Systems (2/2)

❑ A directory structure (per file system) organizes the files

- In UFS, this includes file names and associated inode numbers
- In NTFS, it is stored in the master file table

❑ A per-file FCB contains many details about the file


- It has a unique identifier number to allow association with a directory entry
- In NTFS, this information is actually stored within the master file table, which uses a relational database structure, with a row per file

In-Memory File-System Structures

- ❑ An in-memory **mount table** contains information about each mounted volume
- ❑ An in-memory directory-structure cache holds the directory information of recently accessed directories
- ❑ The **system-wide open-file table** contains a copy of the FCB of each open file
 - As well as other information
- ❑ The **per-process open-file table** contains pointers to the appropriate entries in the system-wide open-file table
 - As well as other information
- ❑ Buffers hold file-system blocks when they are being read from or written to a file system

File Open and File Read

- ❑ `open ()` returns a file handle (descriptor) for subsequent use



Outline

- ❑ File-System Structure
- ❑ File-System Operations
- ❑ **Directory Implementation**
- ❑ Allocation Methods
- ❑ Free-Space Management
- ❑ Efficiency and Performance
- ❑ Recovery
- ❑ Example: WAFL File System

Directory Implementation

❑ Linear list of file names with pointers to the data blocks

- Simple to program
- Time-consuming (linear search) to execute
 - A sorted list allows a binary search and decreases the average search time

❑ Hash table: linear list with a hash data structure

- Decrease the directory search time
- Need some provision for collisions
 - Situations where two file names hash to the same location
- Have difficulties with its generally fixed size and the dependence of the hash function on that size
 - Alternatively, we can use a chained-overflow hash table

Outline

- ❑ File-System Structure
- ❑ File-System Operations
- ❑ Directory Implementation
- ❑ **Allocation Methods**
 - Contiguous Allocation
 - Linked Allocation
 - Indexed Allocation
 - Performance
- ❑ Free-Space Management
- ❑ Efficiency and Performance
- ❑ Recovery
- ❑ Example: WAFL File System

Allocation Methods

- ❑ How to allocate space to these files so that storage space is utilized effectively and files can be accessed quickly

Outline

- ❑ File-System Structure
- ❑ File-System Operations
- ❑ Directory Implementation
- ❑ **Allocation Methods**
 - **Contiguous Allocation**
 - Linked Allocation
 - Indexed Allocation
 - Performance
- ❑ Free-Space Management
- ❑ Efficiency and Performance
- ❑ Recovery
- ❑ Example: WAFL File System


Contiguous Allocation

❑ Each file occupy a set of contiguous blocks on the device

- Easy implementation
- External fragmentation
 - Dynamic storage-allocation problem (Section 9.2)
 - Free-space management (Section 14.5)
- One strategy: copying an entire file system onto another device and then coping back
 - Compaction
 - Off-line (during down time)
 - On-line

Directory

file	start	length
count	0	2
tr	14	3
mail	19	6
list	28	4
f	6	2



Extent-Based Systems

- ❑ How much space is needed for a file?
- ❑ How to make the file larger in place?
 - Terminate the program, allocate more space, and run the program again
 - These repeated runs may be costly
 - To prevent them, the user will normally overestimate the amount of space needed
 - Find a larger hole, copy the contents of the file to the new space, and release the previous space
- ❑ If that amount proves not to be large enough, another chunk of contiguous space, known as an extent, is added
 - The location of a file's blocks is then recorded as a location and a block count, plus a link to the first block of the next extent

Outline

- ❑ File-System Structure
- ❑ File-System Operations
- ❑ Directory Implementation
- ❑ **Allocation Methods**
 - Contiguous Allocation
 - **Linked Allocation**
 - Indexed Allocation
 - Performance
- ❑ Free-Space Management
- ❑ Efficiency and Performance
- ❑ Recovery
- ❑ Example: WAFL File System

Linked Allocation

❑ Each file is a linked list of storage blocks which may be scattered anywhere on the device


- The directory contains a pointer to the first and last blocks of the file
- Each block contains a pointer to next block

❑ Advantages

- No external fragmentation
- No compaction needed

❑ Disadvantages

- Inefficiency of finding the i-th block of a file
 - **Solution: clusters**
- More space for pointers
- Reliability



File-Allocation Table (FAT)

- ❑ A section of storage at the beginning of each volume is set aside to contain the table
 - The table has one entry for each block and is indexed by block number

Directory Entry




- ❑ **Advantage**

- Less direct-access time

- ❑ **Disadvantage**

- Probably more disk head seeks



Outline

- ❑ File-System Structure
- ❑ File-System Operations
- ❑ Directory Implementation
- ❑ **Allocation Methods**
 - Contiguous Allocation
 - Linked Allocation
 - **Indexed Allocation**
 - Performance
- ❑ Free-Space Management
- ❑ Efficiency and Performance
- ❑ Recovery
- ❑ Example: WAFL File System

Indexed Allocation

❑ Bring all the pointers together into the index block

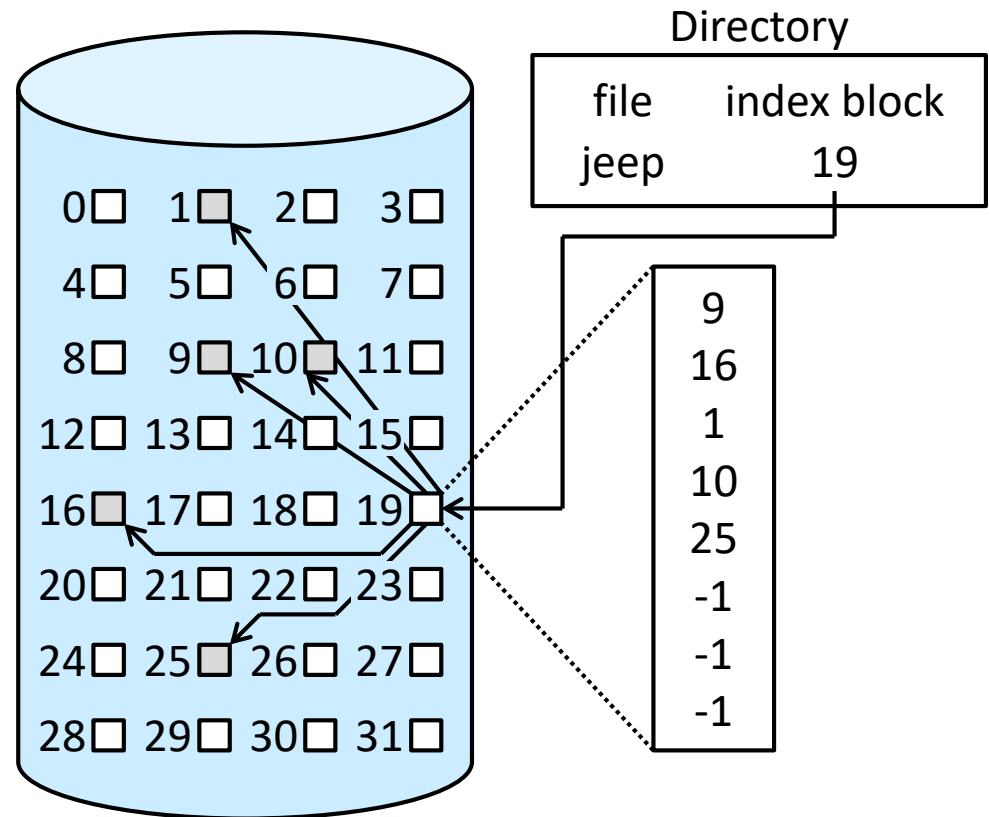
- In the absence of a FAT, linked allocation cannot support efficient direct access

❑ Advantages

- No external fragmentation
- No compaction needed

❑ Disadvantage

- More space than linked allocation



Schemes of Index Blocks (1/2)

❑ Linked scheme

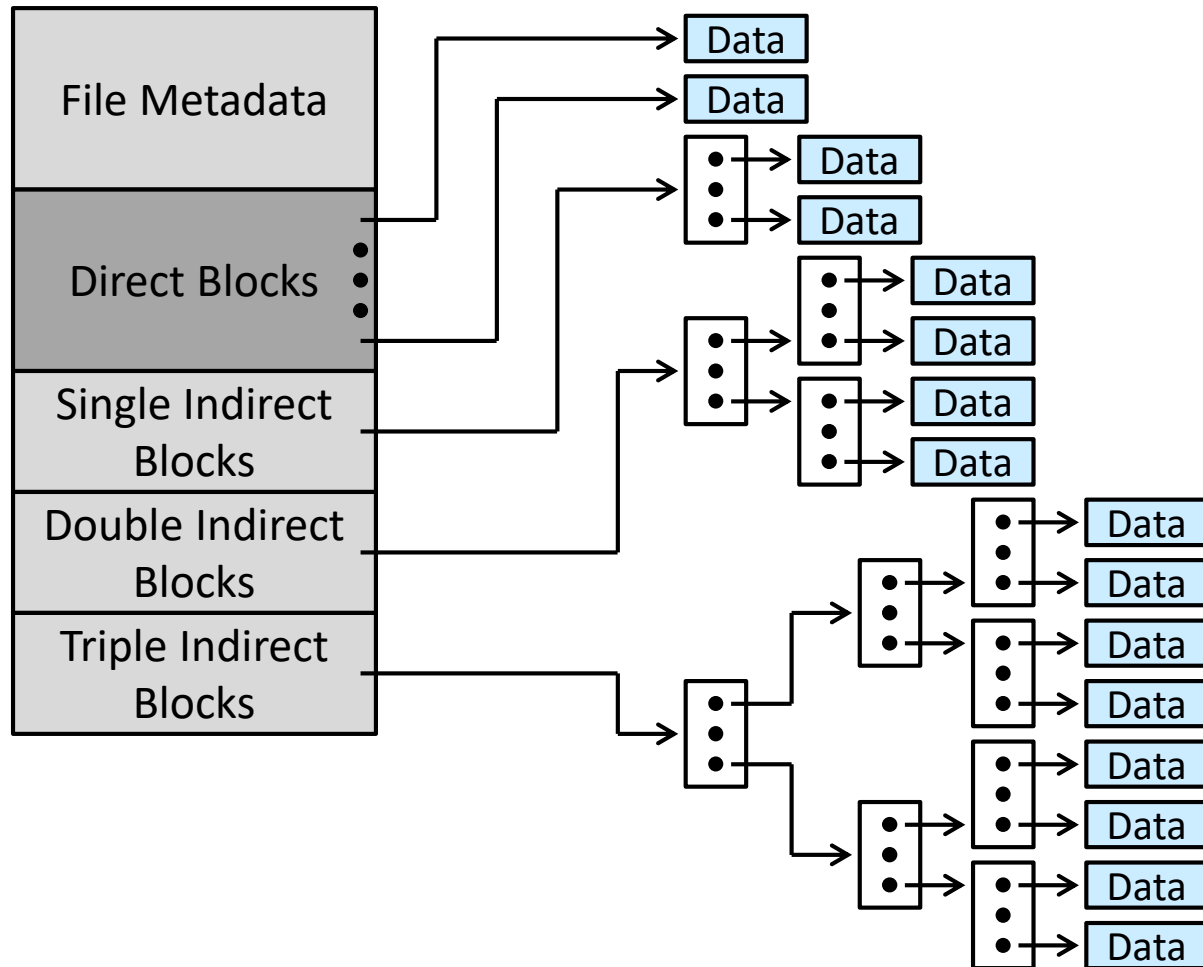
- An index block is normally one storage block
- To allow for large files, we can link together several index blocks

❑ Multilevel index

- A first-level index block points to a set of second-level index blocks, which in turn point to the file blocks
- This approach could be continued to a third or fourth level, depending on the desired maximum file size
- Example
 - With 4,096-byte blocks, we could store 1,024 four-byte pointers in an index block
 - Two levels of indexes allow 1,048,576 data blocks and a file size of up to 4 GB

Schemes of Index Blocks (2/2)

❑ Combined scheme (in UNIX-based file systems)



Outline

- ❑ File-System Structure
- ❑ File-System Operations
- ❑ Directory Implementation
- ❑ **Allocation Methods**
 - Contiguous Allocation
 - Linked Allocation
 - Indexed Allocation
 - **Performance**
- ❑ Free-Space Management
- ❑ Efficiency and Performance
- ❑ Recovery
- ❑ Example: WAFL File System

Performance

- ❑ Goals: storage efficiency and data-block access times
- ❑ A system with mostly sequential access should not use the same method as a system with mostly direct (random) access
 - Contiguous allocation requires only one access to get a block, for any type of access
 - Linked allocation should not be used for direct access
- ❑ Indexed allocation is more complex
 - Depend on the index structure, the file size, and the block position
- ❑ Mix of contiguous, linked, and indexed allocations
- ❑ For NVM, different algorithms and optimizations are needed
 - Reduce the instruction count and overall path between the storage device and application access to the data

Outline

- ❑ File-System Structure
- ❑ File-System Operations
- ❑ Directory Implementation
- ❑ Allocation Methods
- ❑ **Free-Space Management**
- ❑ Efficiency and Performance
- ❑ Recovery
- ❑ Example: WAFL File System

Q&A